

# ECE 657A - Assignment 2

Date Submitted: Mar 09, 2022

## Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import random
from pprint import pprint
from functools import reduce
from scipy import stats
from sklearn import neighbors
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
from sklearn.metrics import
accuracy_score, confusion_matrix, ConfusionMatrixDisplay, classification_report, ma
precision_score, recall_score, f1_score
from sklearn.decomposition import PCA
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.manifold import TSNE
from sklearn.model_selection import
cross_val_score, KFold, StratifiedKFold, GridSearchCV, cross_validate
from sklearn.naive_bayes import MultinomialNB, ComplementNB, GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
import graphviz
from sklearn.ensemble import
RandomForestClassifier, GradientBoostingClassifier, AdaBoostClassifier

import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style="ticks", color_codes=True)
import warnings
warnings.filterwarnings("ignore")
```

## Wine Dataset

## Preprocessing

## Loading Data

```

In [2]: D = ['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density', 'pH',
'sulphates', 'alcohol']
L = 'quality'
C = 'color'
DL = D + [L]
DC = D + [C]
DLC = DL + [C]
#Loading Data set
wine_r = pd.read_csv("winequality-red.csv", sep=';')
#Loading Data set
wine_w = pd.read_csv("winequality-white.csv", sep=';')
wine_w = wine_w.copy()
wine_w[C] = np.zeros(wine_w.shape[0])
wine_r[C] = np.ones(wine_r.shape[0])
wine = pd.concat([wine_w, wine_r])

wine['color'] = wine['color'].astype('category')
wine['quality'] = wine['quality'].astype('category')

```

## Normalization and Scaling

```

In [3]: #Z-Score Normalization
wine_z = wine.copy()
wine_z = wine_z.drop(columns = ['color' , 'quality'])
wine_znormalized = wine_z.apply(stats.zscore)
wine_znormalized['quality'] = wine['quality']
wine_znormalized['color'] = wine['color']

target_z = wine_znormalized['quality']
data_z = wine_znormalized.drop(columns = 'quality')

X_train_z, X_test_z, Y_train_z, Y_test_z = train_test_split(data_z, target_z,
test_size=0.2 , random_state=27)

```

```

In [4]: #MinMax Scaling
wine_s = wine.copy()
wine_s = wine_s.drop(columns = ['color' , 'quality'])
wine_minmax = (wine_s - wine_s.min()) / ( wine_s.max() - wine_s.min())
wine_minmax['quality'] = wine['quality']
wine_minmax['color'] = wine['color']

wine_min_max = wine_minmax
target = wine_min_max['quality']

```

```
data = wine_min_max.drop(columns = 'quality')
X_train, X_test, Y_train, Y_test = train_test_split(data, target,
test_size=0.2 , random_state=27)
```

## Best formulated KNN from Assignment-1

```
In [5]: knn = KNeighborsClassifier(n_neighbors=23 , metric = 'manhattan' , weights =
"distance")
knn.fit(X_train, Y_train)
Y_pred = knn.predict(X_test)

knn_train_acc = knn.score(X_train, Y_train)
print("Training Score: ", knn_train_acc)
knn_test_acc = knn.score(X_test, Y_test)
print("Testing Score: ", knn_test_acc)
```

Training Score: 1.0

Testing Score: 0.68

## Representation Learning

### Independent and dependent variable

```
In [6]: #Scaled 'data' and 'target'
x=data
y=target
n_components=2
target_names = np.sort(y.unique())
```

## PCA

```
In [7]: print("PCA Model - Explained Variance of each component")
pca = PCA(n_components=n_components)
X_r = pca.fit(x).transform(x)
exp_var_ratio_df = pd.DataFrame(pca.explained_variance_ratio_, columns=
['Explained Variance Ratio'])
print(exp_var_ratio_df)

plt.figure(figsize=(7,7))
#colors = plt.cm.get_cmap("twilight")


for i, target_name in enumerate(target_names):
    plt.scatter(
        X_r[y == i, 0], X_r[y == i, 1], alpha=0.6, lw=1,
        label=target_name, cmap="twilight"
    )
```

```
plt.ylim(-0.5,0.7)
plt.legend(loc="best", shadow=False, scatterpoints=1)
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.title("First Two Components from PCA of wine dataset")
#plt.set_cmap("hsv")

plt.show()
```

PCA Model - Explained Variance of each component

	Explained Variance Ratio
0	0.692196
1	0.118142



## LDA

In [8]:

```
print("LDA Model - Explained Variance of each component")
lda = LinearDiscriminantAnalysis(n_components=n_components)
X_r2 = lda.fit(x, y).transform(x)
exp_var_ratio_df = pd.DataFrame(lda.explained_variance_ratio_, columns=
['Explained Variance Ratio'])
print(exp_var_ratio_df)

plt.figure(figsize=(7,7))
#colors = plt.cm.get_cmap("twilight")
for i, target_name in enumerate(target_names):
    plt.scatter(
```

```

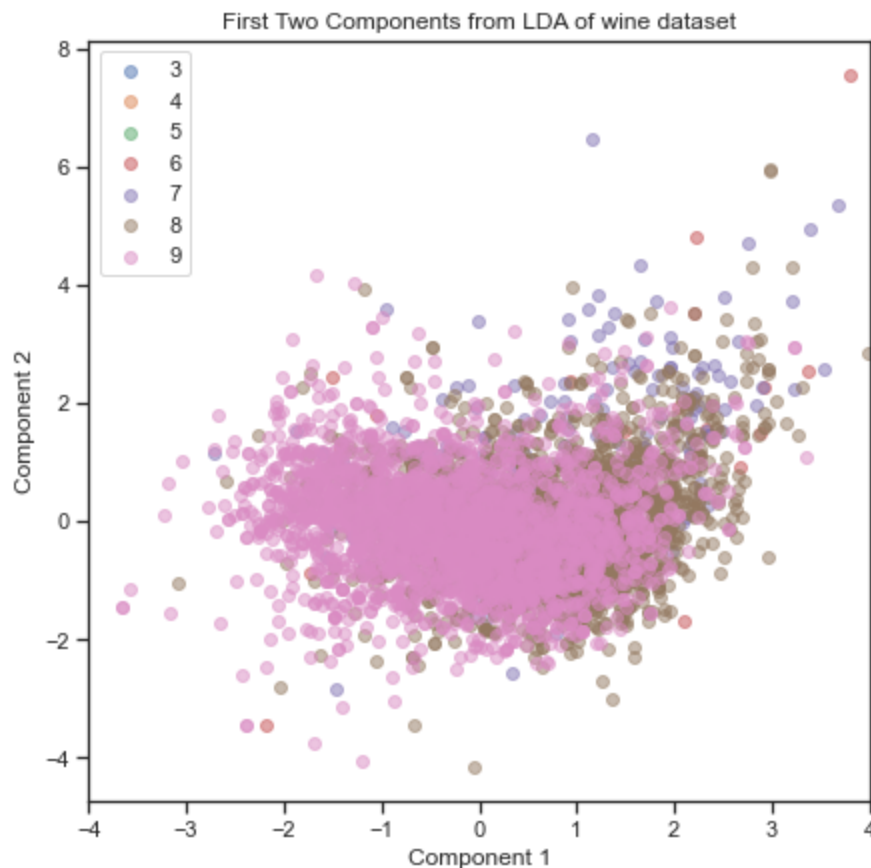
        X_r2[y == i, 0], X_r2[y == i, 1], alpha=0.5, lw=1, label=target_name,
        cmap="twilight"
    )
plt.xlim(-4,4)
plt.legend(loc="best", shadow=False, scatterpoints=1)
plt.xlabel("Component 1")
plt.ylabel("Component 2")
plt.title("First Two Components from LDA of wine dataset")
#plt.set_cmap("twilight")

plt.show()

```

LDA Model - Explained Variance of each component

	Explained Variance Ratio
0	0.861621
1	0.088287



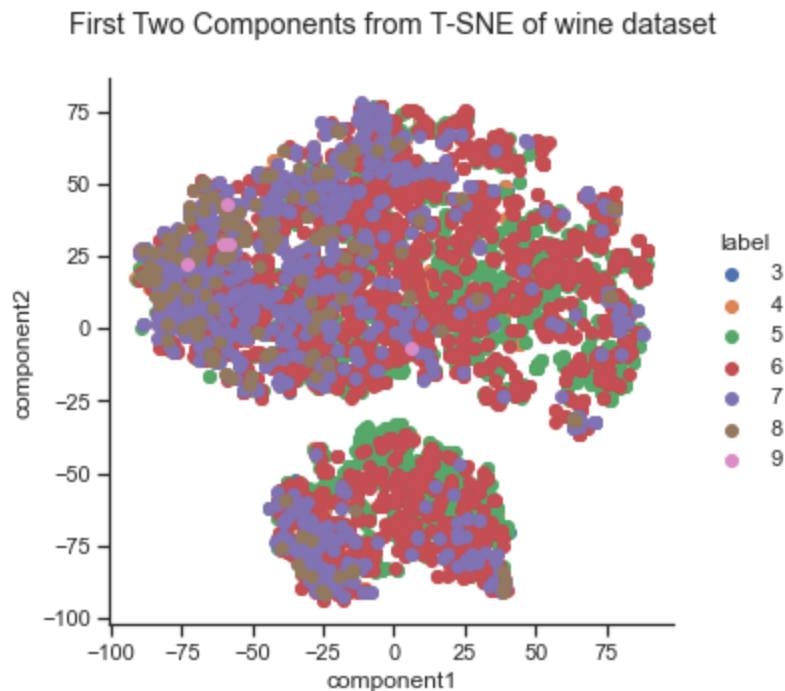
## T-SNE

```

In [9]: tsne = TSNE(n_components=n_components)
X_r3 = tsne.fit_transform(x)
tsne_data = np.vstack((X_r3.T, target)).T
tsne_df = pd.DataFrame(data = tsne_data, columns = ['component1',
'component2', 'label'] )
#print(tsne_df.head())
tsne_df['label'] = tsne_df['label'].astype('int')
tsne_df['label'] = tsne_df['label'].astype('category')

```

```
g = sns.FacetGrid(tsne_df, hue='label', height=5).map(plt.scatter,
'component1', 'component2')
g.add_legend()
g.fig.suptitle("First Two Components from T-SNE of wine dataset")
g.tight_layout()
```




## Comments

- PCA: Separation on X-axis (PC1) is more significant and two clusters are formed, spread among each cluster along X-axis is less than Y-axis (color is categorical variable and could be the reason why two different clusters are formed), we have not removed outliers and hence does not perform well.
- LDA: Same class labels are clustered together and is able to better classify between different classes (number of datapoints in various classes are highly different)
- t-SNE: used default settings, non-linear function, clusters similar classes as distance between them are less in both the clusters

## Scree Plot (PCA)

```
In [10]: pca = PCA(n_components=None)
X_r = pca.fit(x).transform(x)
explained_variance_pca = pca.explained_variance_ratio_
fig = plt.figure(figsize=(7,5))
tick_label = ['PC' + str(i) for i in range(1,len(explained_variance_pca)+1)]
plt.bar(range(len(explained_variance_pca)), explained_variance_pca, color =
'c', alpha=0.5, align='center', tick_label=tick_label)
plt.plot(range(len(explained_variance_pca)), explained_variance_pca, color =
'red', marker='o', mew=4, mfc='yellow', ls='--',
linewidth=1.5, label='Explained Variance')
plt.plot(range(len(explained_variance_pca)), explained_variance_pca.cumsum(),
color='m', marker='^', mew=4, mfc='red', ls=':')
```

```
,linewidth=1.5,label='Cummulative Explained Variance')
plt.ylabel('Proportion of variation explained',fontsize = 15)
plt.xlabel('Principal Components',fontsize = 15)
plt.title('Scree plot (PCA)',fontsize = 18)
plt.legend(loc='best',fontsize = 15)
#plt.axis("off")
plt.tight_layout()
plt.show()
```



## Reduced Dimensionality

### PCA

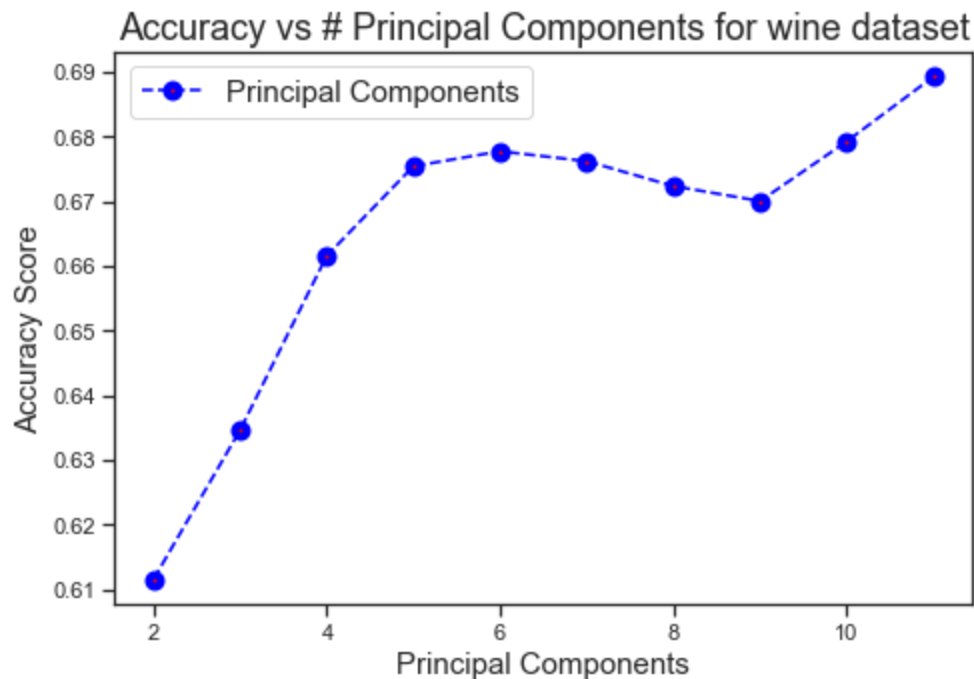
```
In [11]: knn = KNeighborsClassifier(n_neighbors=23 , metric = 'manhattan' , weights =
"distance")
D = data.shape[1]
#print(D)
acc_list = []

for i in range(2,D):
    pca = PCA(n_components=i)
    X_train_r = pca.fit(X_train).transform(X_train)
    X_test_r = pca.transform(X_test)
    knn.fit(X_train_r, Y_train)
    y_pred = knn.predict(X_test_r)
    acc_list.append(metrics.accuracy_score(Y_test,y_pred))

fig = plt.figure(figsize=(7,5))
plt.plot(range(2,D),acc_list, color = 'blue',marker = 'o',mew =4,mfc='red',ls
='--' ,linewidth=1.5,label='Principal Components')
```

```
plt.ylabel('Accuracy Score',fontsize = 15)
plt.xlabel('Principal Components',fontsize = 15)
plt.title('Accuracy vs # Principal Components for wine dataset',fontsize = 18)
plt.legend(loc='best',fontsize = 15)
plt.tight_layout()
plt.show()

print("The maximum accuracy {} happens to be when number of Principal Components is : {}".format(max(acc_list),acc_list.index(max(acc_list))+2))
```



The maximum accuracy 0.6892307692307692 happens to be when number of Principal Components is : 11

## LDA

```
In [12]: knn = KNeighborsClassifier(n_neighbors=23 , metric = 'manhattan' , weights = "distance")


#LDA is constrained by the number of components it can use , given by :
min(n_classes-1, len(features))
D = min(data.shape[1],len(target_names))
acc_list =[]

for i in range(2,D):
    lda = LinearDiscriminantAnalysis(n_components=i)
    X_train_r = lda.fit(X_train,Y_train).transform(X_train)
    X_test_r = lda.transform(X_test)
    knn.fit(X_train_r, Y_train)
    y_pred = knn.predict(X_test_r)
    acc_list.append(metrics.accuracy_score(Y_test,y_pred))
```



```
fig = plt.figure(figsize=(7,5))
plt.plot(range(2,D),acc_list, color = 'blue',marker = 'o',mew =4,mfc='red',ls
='--' ,linewidth=1.5,label='LDA Components')
plt.ylabel('Accuracy Score',fontsize = 12)
plt.xlabel('LDA Components',fontsize = 12)
plt.title('Accuracy vs # of LDA Components for wine dataset',fontsize = 16)
plt.legend(loc='best',fontsize = 12)
plt.tight_layout()
plt.show()

print("The maximum accuracy {} happens to be when number of LDA Components is
: {}".format(max(acc_list),acc_list.index(max(acc_list))+2))
```



The maximum accuracy 0.676923076923077 happens to be when number of LDA Components is : 5

## Comments

- Have used 5 principal components in subsequent parts as there is not a significant difference when included more components. Also, 11 components is the original dimension of data and hence would be better representation but there is no significant loss of accuracy by including 5 components.
- LDA gives better accuracy with 5 components.
- Accuracy with reduced dimensions is comparable with accuracy from Assignment-1 (0.68). PCA performs slightly better with 0.69 accuracy. The major observation is that with reduced dimensions accuracy is comparable with the original dataset (with all features) hence PCA and dLDA are able to capture most information in lesser components.

## Versions of dataset (raw/pca/lda)

```
In [7]: wine_raw = wine.copy()
wine_normalized_raw = wine_znormalized.copy()
wine_minmax_normalized = wine_minmax.copy()
```

```

pca = PCA(n_components=5)
X_r = pca.fit(x).transform(x)
pca_data = np.vstack((X_r.T, target)).T
wine_pca = pd.DataFrame(data = pca_data, columns =
['component1', 'component2', 'component3', 'component4', 'component5', 'label'] )
wine_pca_temp = wine_pca.copy()
wine_pca_temp_2 = wine_pca_temp.drop(columns = ['label'])
wine_pca = (wine_pca_temp_2 - wine_pca_temp_2.min(axis=0)) / (
wine_pca_temp_2.max(axis=0) - wine_pca_temp_2.min(axis=0))
wine_pca['label']=wine_pca_temp['label']
wine_pca['label'] = wine_pca['label'].astype('int')
wine_pca['label'] = wine_pca['label'].astype('category')
#print(wine_pca.head())

lda = LinearDiscriminantAnalysis(n_components=5)
X_r2 = lda.fit(x, y).transform(x)
lda_data = np.vstack((X_r2.T, target)).T
wine_lda = pd.DataFrame(data = lda_data, columns = ['component1',
'component2', 'component3', 'component4', 'component5', 'label'] )
wine_lda_temp = wine_lda.copy()
wine_lda_temp_2 = wine_lda_temp.drop(columns = ['label'])
wine_lda = (wine_lda_temp_2 - wine_lda_temp_2.min(axis=0)) / (
wine_lda_temp_2.max(axis=0) - wine_lda_temp_2.min(axis=0))
wine_lda['label']=wine_lda_temp['label']
wine_lda['label'] = wine_lda['label'].astype('int')
wine_lda['label'] = wine_lda['label'].astype('category')
#print(wine_lda.head())

```

In [14]:

```

wine_p = wine_pca.copy()
target_pca = wine_p['label']
data_pca = wine_p.drop(columns = 'label')
X_train_pca, X_test_pca, Y_train_pca, Y_test_pca = train_test_split(data_pca,
target_pca, test_size=0.2 , random_state=27)

wine_l = wine_lda.copy()
target_lda = wine_l['label']
data_lda = wine_l.drop(columns = 'label')
X_train_lda, X_test_lda, Y_train_lda, Y_test_lda = train_test_split(data_lda,
target_lda, test_size=0.2 , random_state=27)

```

## Naive Bayes Classifier

### 5-Fold CV for Naive Bayes & KNN classifier

```
In [15]: kfold = KFold(n_splits=5)
         clrs = []

         clrs.append(KNeighborsClassifier(n_neighbors=23 , metric = 'manhattan' ,
         weights = "distance"))
         clrs.append(GaussianNB())
         clrs.append(MultinomialNB())
         clrs.append(ComplementNB())

         cv_results = []
         cv_results_pca = []
         cv_results_lda = []

         for clr in clrs :
             cv_results.append(cross_val_score(clr, X_train, Y_train , scoring =
             'accuracy', cv = kfold, n_jobs=-1))
             cv_results_pca.append(cross_val_score(clr, X_train_pca, Y_train_pca ,
             scoring = 'accuracy', cv = kfold, n_jobs=-1))
             cv_results_lda.append(cross_val_score(clr, X_train_lda, Y_train_lda ,
             scoring = 'accuracy', cv = kfold, n_jobs=-1))

         cv_means = []
         cv_means_pca = []
         cv_means_lda = []

         for cv_result in cv_results:
             cv_means.append(cv_result.mean())

         for cv_result in cv_results_pca:
             cv_means_pca.append(cv_result.mean())

         for cv_result in cv_results_lda:
             cv_means_lda.append(cv_result.mean())

         cv_df = pd.DataFrame({"Mean_Accuracy":cv_means,"Algo":
         ['KNeighborsClassifier','GaussianNB','MultinomialNB','ComplementNB']})
         cv_df_pca = pd.DataFrame({"Mean_Accuracy":cv_means_pca,"Algo":
         ['KNeighborsClassifier','GaussianNB','MultinomialNB','ComplementNB']})
         cv_df_lda = pd.DataFrame({"Mean_Accuracy":cv_means_lda,"Algo":
         ['KNeighborsClassifier','GaussianNB','MultinomialNB','ComplementNB']})

         #print(cv_df)
         #print(cv_df_pca)
         #print(cv_df_lda)
```

In [16]:

```
data_frames = [cv_df, cv_df_pca, cv_df_lda]
df_merged = reduce(lambda left, right: pd.merge(left, right, on=
['Algo'], how='outer'), data_frames)
df_merged.rename(columns={"Mean_Accuracy_x": "wine_raw", "Mean_Accuracy_y":
"wine_pca", "Mean_Accuracy": "wine_lda"}, inplace=True)
df_merged.set_index('Algo', inplace=True)
print(df_merged)
```

	wine_raw	wine_pca	wine_lda
Algo			
KNeighborsClassifier	0.652869	0.640555	0.646522
GaussianNB	0.359820	0.509904	0.537423
MultinomialNB	0.472769	0.436018	0.436018
ComplementNB	0.396376	0.367124	0.438711

In [17]:

```
kfold = KFold(n_splits=5)
clrs = []

clrs.append(KNeighborsClassifier(n_neighbors=23 , metric = 'manhattan' ,
weights = "distance"))
clrs.append(GaussianNB())
clrs.append(MultinomialNB())
clrs.append(ComplementNB())

cv_results = []
cv_results_pca = []
cv_results_lda = []

#scoring = {'accuracy' : make_scorer(accuracy_score),
#           'precision' : make_scorer(precision_score),
#           'recall' : make_scorer(recall_score),
#           'f1_score' : make_scorer(f1_score)}

for clr in clrs :
    cv_results.append(cross_val_score(clr, data, target , scoring =
'accuracy', cv = kfold, n_jobs=-1))
    cv_results_pca.append(cross_val_score(clr, data_pca, target_pca , scoring
= 'accuracy', cv = kfold, n_jobs=-1))
    cv_results_lda.append(cross_val_score(clr, data_pca, target_pca , scoring
= 'accuracy', cv = kfold, n_jobs=-1))

cv_means = []
cv_means_pca = []
cv_means_lda = []
```

```

#print(cv_results)
for cv_result in cv_results:
    cv_means.append(cv_result.mean())

for cv_result in cv_results_pca:
    cv_means_pca.append(cv_result.mean())

for cv_result in cv_results_lda:
    cv_means_lda.append(cv_result.mean())

cv_df = pd.DataFrame({"Mean_Accuracy":cv_means,"Algo":
['KNeighborsClassifier','GaussianNB','MultinomialNB','ComplementNB']})
cv_df_pca = pd.DataFrame({"Mean_Accuracy":cv_means_pca,"Algo":
['KNeighborsClassifier','GaussianNB','MultinomialNB','ComplementNB']})
cv_df_lda = pd.DataFrame({"Mean_Accuracy":cv_means_lda,"Algo":
['KNeighborsClassifier','GaussianNB','MultinomialNB','ComplementNB']})

#print(cv_df)
#print(cv_df_pca)
#print(cv_df_lda)

```

In [18]:

```

data_frames = [cv_df, cv_df_pca, cv_df_lda]
df_merged = reduce(lambda left,right: pd.merge(left,right,on=
['Algo'],how='outer'), data_frames)
df_merged.rename(columns={"Mean_Accuracy_x": "wine_raw", "Mean_Accuracy_y":
"wine_pca","Mean_Accuracy":"wine_lda"},inplace=True)
df_merged.set_index('Algo',inplace=True)
print(df_merged)

```

	wine_raw	wine_pca	wine_lda
Algo			
KNeighborsClassifier	0.504236	0.489310	0.489310
GaussianNB	0.309505	0.467139	0.467139
MultinomialNB	0.430356	0.436515	0.436515
ComplementNB	0.396336	0.385406	0.385406

## Comments

- Complement Naive Bayes does not outperform as expected for the unbalance wine dataset. While training using Complement NB it trains each class with all data, but the sample from that class and because it is a multiclass classification with higher Prior, the likelihood impact is reduced and overall accuracy is not good.

## Decision Trees Classifier

### 5-Fold CV and hyperparameter tuning

In [19]:

```

dtc = DecisionTreeClassifier()
dtc.fit(X_train, Y_train)
Y_pred = dtc.predict(X_test)
#cm=confusion_matrix(Y_test, Y_pred)
#print(confusion_matrix(Y_test, Y_pred))
#print(classification_report(Y_test, Y_pred))
print("Accuracy Score for Decision Tree classifier without hyper parameter
tuning is: ",accuracy_score(Y_test, Y_pred))
#disp =
ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=dtc.classes_)
#disp.plot()

grid_params = {
    'criterion' : ['gini', 'entropy'],
    'max_depth' : range(2,25,1),
    'min_samples_split' : range(2, 10, 1),
    'min_samples_leaf' : range(2, 10, 1)
}

grid_search = GridSearchCV(dtc, grid_params, cv = 5, n_jobs = -1, verbose =
1)
grid_search.fit(X_train, Y_train)
dtc = grid_search.best_estimator_
Y_pred = dtc.predict(X_test)
print("Accuracy Score is :",accuracy_score(Y_test, Y_pred))
# best parameters and best score
print("Best Parameters are : ",grid_search.best_params_)
print("Best Score is : ",grid_search.best_score_)
#dt_raw_acc = accuracy_score(Y_test, Y_pred)
dt_raw_set=grid_search.best_params_

grid_search.fit(X_train_pca, Y_train_pca)
dtc = grid_search.best_estimator_
Y_pred_pca = dtc.predict(X_test_pca)
print("Accuracy Score using PCA is :",accuracy_score(Y_test_pca, Y_pred_pca))
# best parameters and best score
print("Best Parameters using PCA are : ",grid_search.best_params_)
print("Best Score using PCA is : ",grid_search.best_score_)
dt_pca_set=grid_search.best_params_

grid_search.fit(X_train_lda, Y_train_lda)
dtc = grid_search.best_estimator_
Y_pred_lda = dtc.predict(X_test_lda)
print("Accuracy Score using LDA is :",accuracy_score(Y_test_lda, Y_pred_lda))

```

```
# best parameters and best score
print("Best Parameters using LDA are : ",grid_search.best_params_)
print("Best Score using LDA is : ",grid_search.best_score_)
dt_lda_set=grid_search.best_params_
```

Accuracy Score for Decision Tree classifier without hyper parameter tuning is: 0.6115384615384616

Fitting 5 folds for each of 2944 candidates, totalling 14720 fits

Accuracy Score is : 0.5792307692307692

Best Parameters are : {'criterion': 'entropy', 'max\_depth': 18, 'min\_samples\_leaf': 2, 'min\_samples\_split': 6}

Best Score is : 0.5639766417413193

Fitting 5 folds for each of 2944 candidates, totalling 14720 fits

Accuracy Score using PCA is : 0.5584615384615385

Best Parameters using PCA are : {'criterion': 'entropy', 'max\_depth': 18, 'min\_samples\_leaf': 2, 'min\_samples\_split': 2}

Best Score using PCA is : 0.5474289257422077

Fitting 5 folds for each of 2944 candidates, totalling 14720 fits

Accuracy Score using LDA is : 0.5423076923076923

Best Parameters using LDA are : {'criterion': 'gini', 'max\_depth': 8, 'min\_samples\_leaf': 2, 'min\_samples\_split': 6}

Best Score using LDA is : 0.5505084400681128

In [20]:

```
dt_raw_acc = accuracy_score(Y_test, Y_pred)
dt_pca_acc = accuracy_score(Y_test_pca, Y_pred_pca)
dt_lda_acc = accuracy_score(Y_test_lda, Y_pred_lda)
#print(dt_raw_acc,dt_pca_acc,dt_lda_acc)
```

In [21]:

```
cm=confusion_matrix(Y_test, Y_pred)
cm_pca=confusion_matrix(Y_test_pca, Y_pred_pca)
cm_lda=confusion_matrix(Y_test_lda, Y_pred_lda)
cm_dict = {
    "raw": cm,
    "pca": cm_pca,
    "lda": cm_lda
}


f, axes = plt.subplots(1, 3, figsize=(20, 5), sharey='row')
for i,(c,cm) in enumerate(cm_dict.items()):
    #print(i,cm)
    disp = ConfusionMatrixDisplay(cm,display_labels=dtc.classes_)
    disp.plot(ax=axes[i], xticks_rotation=45)
    disp.ax_.set_title(c)
    disp.im_.colorbar.remove()
    disp.ax_.set_xlabel('')
    if i!=0:
        disp.ax_.set_ylabel('')
```

```

f.text(0.4, 0.1, 'Predicted label', ha='left')
plt.subplots_adjust(wspace=0.40, hspace=0.1)

f.suptitle('Confusion matrix', fontsize=12)
f.colorbar(dispatch.im_, ax=axes)
plt.show()

```



## Accuracy vs Max\_depth

In [22]:

```

def plot_acc_vs_depth(dataset):
    max_depth_range = list(range(1, 25))
    acc_list = []
    acc_list_pca = []
    acc_list_lda = []
    for depth in max_depth_range:
        dtc_w = DecisionTreeClassifier(max_depth = depth)
        if dataset[0] == 'raw':
            dtc_w.fit(X_train, Y_train)
            score = dtc_w.score(X_test, Y_test)
            acc_list.append(score)
        if dataset[1] == 'pca':
            dtc_w.fit(X_train_pca, Y_train_pca)
            score = dtc_w.score(X_test_pca, Y_test_pca)
            acc_list_pca.append(score)
        if dataset[2] == 'lda':
            dtc_w.fit(X_train_lda, Y_train_lda)
            score = dtc_w.score(X_test_lda, Y_test_lda)
            acc_list_lda.append(score)

    return (acc_list, acc_list_pca, acc_list_lda)

acc_list, acc_list_pca, acc_list_lda = plot_acc_vs_depth(['raw', 'pca', 'lda'])
fig = plt.figure(figsize=(7, 5))
plt.plot(range(1, 25), acc_list, color = 'blue', marker = 'o', mew

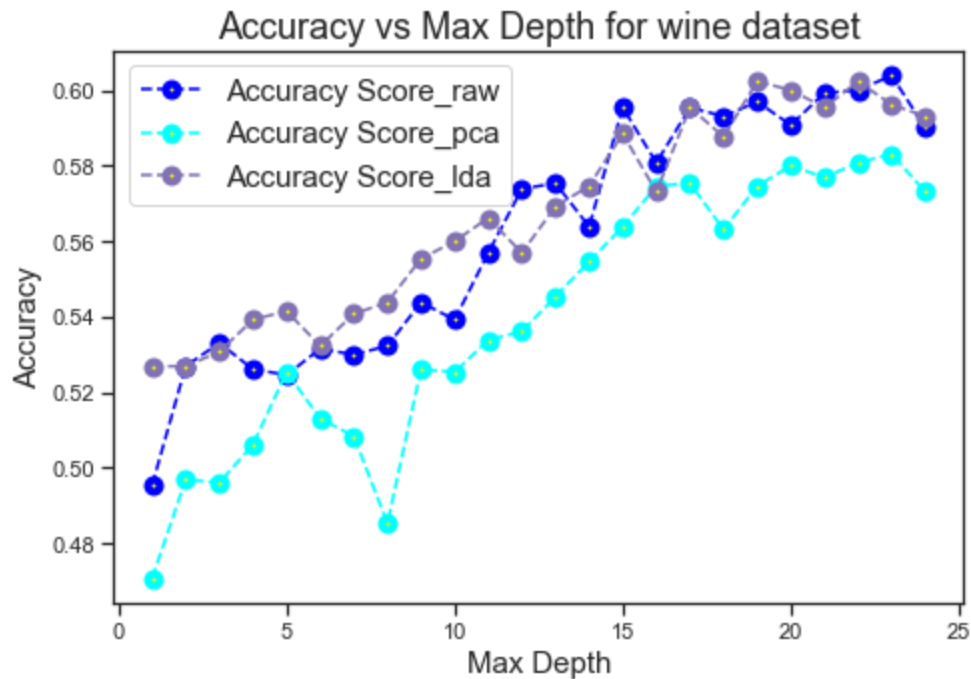
```



```

=4,mfc='yellow',ls='--',linewidth=1.5,label='Accuracy Score_raw')
plt.plot(range(1,25),acc_list_pca, color = 'cyan',marker='o',mew
=4,mfc='yellow',ls='--',linewidth=1.5,label='Accuracy Score_pca')
plt.plot(range(1,25),acc_list_lda, color = 'm',marker='o',mew
=4,mfc='yellow',ls='--',linewidth=1.5,label='Accuracy Score_lda')
plt.ylabel('Accuracy',fontsize = 15)
plt.xlabel('Max Depth',fontsize = 15)
plt.title('Accuracy vs Max Depth for wine dataset',fontsize = 18)
plt.legend(loc='best',fontsize = 15)
plt.tight_layout()
plt.show()

```



## Splitting rules used for the trees

```

In [23]: dtc_g = DecisionTreeClassifier(criterion= 'entropy', max_depth=18,
min_samples_leaf= 2, min_samples_split=5)
dtc_g.fit(X_train,Y_train)
fn=X_train.columns
cn = wine_raw.quality.unique().astype(str)
dot_data = tree.export_graphviz(dtc_g,
filled=True,feature_names=fn,class_names=cn ,rounded=True,
special_characters=True)
graph = graphviz.Source(dot_data)
graph

```

Out[23]:



## Text split rules

```

In [8]: from sklearn.tree import export_text

```

```
dtc_g = DecisionTreeClassifier(criterion= 'entropy', max_depth=18,  
min_samples_leaf= 2, min_samples_split=5)  
dt_rules = dtc_g.fit(X_train,Y_train)  
fn_t=X_train.columns  
#print(type(fn_t.tolist()))  
tree_rules = export_text(dt_rules,feature_names = fn_t.tolist())  
#print(tree_rules)
```

#### Splitting rules used for the trees but fitted on page

In [10]:

```
dtc_g = DecisionTreeClassifier(criterion= 'entropy', max_depth=18,  
min_samples_leaf= 2, min_samples_split=5)  
dtc_g.fit(X_train,Y_train)  
fn=X_train.columns  
cn = wine_raw.quality.unique().astype(str)  
fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (8,8), dpi=300)  
tree.plot_tree(dtc_g, filled=True,feature_names=fn,class_names=cn  
,rounded=True)  
fig.savefig('dtc_tree_plot_4.png')
```



```

print("Accuracy Score for Forest classifier without hyper parameter tuning
is: ",accuracy_score(Y_test, Y_pred))
#disp =
ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=dtc.classes_)
#disp.plot()

grid_params = {
    'n_estimators': [2,5,200, 500],
    'criterion' : ['entropy'],
    'max_depth' : range(2,25,1),
    #'max_features': range(3,10,1),
    #'min_samples_split' : range(2, 10, 1),
    #'min_samples_leaf' : range(2, 10, 1)
}

grid_search = GridSearchCV(rfc, grid_params, cv = 5, n_jobs = -1, verbose =
True)
#print(grid_search.get_params().keys())
grid_search.fit(X_train, Y_train)
rfc = grid_search.best_estimator_
Y_pred = rfc.predict(X_test)
print("Accuracy Score is :",accuracy_score(Y_test, Y_pred))
rf_raw_acc=accuracy_score(Y_test, Y_pred)
# best parameters and best score
print("Best Parameters are : ",grid_search.best_params_)
rf_raw_set=grid_search.best_params_
#print("Best Score is : ",grid_search.best_score_)
ac_df=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
ac_df["accuracy"]=grid_search.cv_results_['mean_test_score']

grid_search.fit(X_train_pca, Y_train_pca)
rfc_pca = grid_search.best_estimator_
Y_pred_pca = rfc_pca.predict(X_test_pca)
print("Accuracy Score using PCA is :",accuracy_score(Y_test_pca, Y_pred_pca))
rf_pca_acc=accuracy_score(Y_test_pca, Y_pred_pca)
# best parameters and best score
print("Best Parameters using PCA are : ",grid_search.best_params_)
rf_pca_set=grid_search.best_params_
#print("Best Score using PCA is : ",grid_search.best_score_)
ac_df_pca=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
ac_df_pca["accuracy"]=grid_search.cv_results_['mean_test_score']

```

```

grid_search.fit(X_train_lda, Y_train_lda)
rfc_lda = grid_search.best_estimator_
Y_pred_lda = rfc_lda.predict(X_test_lda)
print("Accuracy Score using LDA is :",accuracy_score(Y_test_lda, Y_pred_lda))
rf_lda_acc=accuracy_score(Y_test_lda, Y_pred_lda)
# best parameters and best score
print("Best Parameters using LDA are : ",grid_search.best_params_)
#print("Best Score using LDA is : ",grid_search.best_score_)
rf_lda_set=grid_search.best_params_
ac_df_lda=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
ac_df_lda["accuracy"]=grid_search.cv_results_['mean_test_score']

```

Accuracy Score for Forest classifier without hyper parameter tuning is: 0.6915384615384615

Fitting 5 folds for each of 92 candidates, totalling 460 fits

Accuracy Score is : 0.6953846153846154

Best Parameters are : {'criterion': 'entropy', 'max\_depth': 18, 'n\_estimators': 500}

Fitting 5 folds for each of 92 candidates, totalling 460 fits

Accuracy Score using PCA is : 0.6607692307692308

Best Parameters using PCA are : {'criterion': 'entropy', 'max\_depth': 16, 'n\_estimators': 200}

Fitting 5 folds for each of 92 candidates, totalling 460 fits

Accuracy Score using LDA is : 0.6784615384615384

Best Parameters using LDA are : {'criterion': 'entropy', 'max\_depth': 17, 'n\_estimators': 500}

In [26]:

```


cm=confusion_matrix(Y_test, Y_pred)
cm_pca=confusion_matrix(Y_test_pca, Y_pred_pca)
cm_lda=confusion_matrix(Y_test_lda, Y_pred_lda)
cm_dict = {
    "raw": cm,
    "pca": cm_pca,
    "lda": cm_lda
}

f, axes = plt.subplots(1, 3, figsize=(20, 5), sharey='row')
for i, (c, cm) in enumerate(cm_dict.items()):
    #print(i, cm)
    disp = ConfusionMatrixDisplay(cm, display_labels=dtc.classes_)
    disp.plot(ax=axes[i], xticks_rotation=45)
    disp.ax_.set_title(c)
    disp.im_.colorbar.remove()
    disp.ax_.set_xlabel('')
    if i!=0:
        disp.ax_.set_ylabel('')

```

```
f.text(0.4, 0.1, 'Predicted label', ha='left')
plt.subplots_adjust(wspace=0.40, hspace=0.1)

f.colorbar(dispatch.im_, ax=axes)
plt.show()
```



## Heatmap of max\_depth vs n\_estimators


In [27]:

```
#Pivoting the dataframe for plotting heat map
ac_df=ac_df.pivot(index='max_depth',columns='n_estimators',values='accuracy')
ac_df_pca=ac_df_pca.pivot(index='max_depth',columns='n_estimators',values='accuracy')
ac_df_lda=ac_df_lda.pivot(index='max_depth',columns='n_estimators',values='accuracy')

#Plotting the graph
fig, ax =plt.subplots(1,3,figsize=(20, 5), sharey='row')

sns.heatmap(ac_df,annot=True, ax=ax[0])
sns.heatmap(ac_df_pca,annot=True ,ax=ax[1])
sns.heatmap(ac_df_lda,annot=True ,ax=ax[2])

plt.subplots_adjust(wspace=0.40, hspace=0.1)
fig.suptitle('Accuracy Heatplot', fontsize=12)
plt.show()
```



# Gradient Boosting Classifier

## 5-Fold CV and hyperparameter tuning

In [28]:

```
gbc = GradientBoostingClassifier()
gbc.fit(X_train, Y_train)
Y_pred = gbc.predict(X_test)
cm=confusion_matrix(Y_test, Y_pred)
#print(confusion_matrix(Y_test, Y_pred))
#print(classification_report(Y_test, Y_pred))
print("Accuracy Score for Gradient Boosting classifier without hyper
parameter tuning is: ",accuracy_score(Y_test, Y_pred))
#disp =
ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=dtc.classes_)
#disp.plot()

grid_params = {
    'n_estimators': [2,120,200,500,800],
    'learning_rate' : [0.01, 0.1,1],
    #'criterion' : [None,'mse','mae'],
    #'max_depth' : range(2,22,2),
    #'max_features': range(3,10,1),
    #'min_samples_split' : range(2, 10, 1),
    #'min_samples_leaf' : range(2, 10, 1)
}

grid_search = GridSearchCV(gbc, grid_params, cv = 5, n_jobs = -1, verbose =
True)
#print(grid_search.get_params().keys())
grid_search.fit(X_train, Y_train)
gbc = grid_search.best_estimator_
Y_pred = gbc.predict(X_test)
print("Accuracy Score is :",accuracy_score(Y_test, Y_pred))
gbc_acc=accuracy_score(Y_test, Y_pred)
# best parameters and best score
print("Best Parameters are : ",grid_search.best_params_)
gbc_set=grid_search.best_params_
#print("Best Score using PCA is : ",grid_search.best_score_)
gbc_ac_df=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
gbc_ac_df["accuracy"]=grid_search.cv_results_['mean_test_score']
```

```

grid_search.fit(X_train_pca, Y_train_pca)
gbc = grid_search.best_estimator_
Y_pred_pca = gbc.predict(X_test_pca)
print("Accuracy Score is :",accuracy_score(Y_test_pca, Y_pred_pca))
gbc_pca_acc=accuracy_score(Y_test_pca, Y_pred_pca)
# best parameters and best score
print("Best Parameters are : ",grid_search.best_params_)
gbc_pca_set=grid_search.best_params_
#print("Best Score using PCA is : ",grid_search.best_score_)
gbc_ac_df_pca=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
gbc_ac_df_pca["accuracy"]=grid_search.cv_results_['mean_test_score']

grid_search.fit(X_train_lda, Y_train_lda)
gbc = grid_search.best_estimator_
Y_pred_lda = gbc.predict(X_test_lda)
print("Accuracy Score is :",accuracy_score(Y_test_lda, Y_pred_lda))
gbc_lda_acc=accuracy_score(Y_test_lda, Y_pred_lda)
# best parameters and best score
print("Best Parameters are : ",grid_search.best_params_)
gbc_lda_set=grid_search.best_params_
#print("Best Score using PCA is : ",grid_search.best_score_)
gbc_ac_df_lda=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
gbc_ac_df_lda["accuracy"]=grid_search.cv_results_['mean_test_score']

```

Accuracy Score for Gradient Boosting classifier without hyper parameter tuning is: 0.6007692307692307

Fitting 5 folds for each of 15 candidates, totalling 75 fits

Accuracy Score is : 0.6346153846153846

Best Parameters are : {'learning\_rate': 0.1, 'n\_estimators': 800}

Fitting 5 folds for each of 15 candidates, totalling 75 fits

Accuracy Score is : 0.6192307692307693

Best Parameters are : {'learning\_rate': 0.1, 'n\_estimators': 800}

Fitting 5 folds for each of 15 candidates, totalling 75 fits

Accuracy Score is : 0.6430769230769231

Best Parameters are : {'learning\_rate': 0.1, 'n\_estimators': 800}

In [29]:

```

cm=confusion_matrix(Y_test, Y_pred)
cm_pca=confusion_matrix(Y_test_pca, Y_pred_pca)
cm_lda=confusion_matrix(Y_test_lda, Y_pred_lda)
cm_dict = {
    "raw": cm,
    "pca": cm_pca,
    "lda": cm_lda
}

```




```

f, axes = plt.subplots(1, 3, figsize=(20, 5), sharey='row')
for i, (c, cm) in enumerate(cm_dict.items()):
    #print(i, cm)
    disp = ConfusionMatrixDisplay(cm, display_labels=dtc.classes_)
    disp.plot(ax=axes[i], xticks_rotation=45)
    disp.ax_.set_title(c)
    disp.im_.colorbar.remove()
    disp.ax_.set_xlabel('')
    if i!=0:
        disp.ax_.set_ylabel('')

f.text(0.4, 0.1, 'Predicted label', ha='left')
plt.subplots_adjust(wspace=0.40, hspace=0.1)

f.colorbar(disp.im_, ax=axes)
plt.show()

```



## Heatplot of learning\_rate and n\_estimators

```

In [30]: #Pivoting the dataframe for plotting heat map
gbc_ac_df=gbc_ac_df.pivot(index='learning_rate',columns='n_estimators',values='...

gbc_ac_df_pca=gbc_ac_df_pca.pivot(index='learning_rate',columns='n_estimators',...


gbc_ac_df_lda=gbc_ac_df_lda.pivot(index='learning_rate',columns='n_estimators',...

#Plotting the graph
fig, ax =plt.subplots(1,3,figsize=(20, 5), sharey='row')

sns.heatmap(gbc_ac_df,annot=True, ax=ax[0])
sns.heatmap(gbc_ac_df_pca,annot=True ,ax=ax[1])
sns.heatmap(gbc_ac_df_lda,annot=True ,ax=ax[2])

```

```
plt.subplots_adjust(wspace=0.40, hspace=0.1)
fig.suptitle('Accuracy Heatplot', fontsize=12)
plt.show()
```



## Final Results

```
In [51]: dt_dict = {'Algo': ['Decision Tree', 'Random Forest', 'Gradient Boosting'],
                  'wine_raw': [dt_raw_acc, rf_raw_acc, gbc_acc], 'wine_pca':
[dt_pca_acc, rf_pca_acc, gbc_pca_acc],
                  'wine_lda': [dt_lda_acc, rf_lda_acc, gbc_lda_acc]}

df_mid=pd.DataFrame(data=dt_dict, columns=
['Algo', 'wine_raw', 'wine_pca', 'wine_lda'])
#display(df_mid)
df_nb=df_merged.reset_index()
final_df = df_nb.append(df_mid, ignore_index=True)
display(final_df)
```

	Algo	wine_raw	wine_pca	wine_lda
0	KNeighborsClassifier	0.504236	0.489310	0.489310
1	GaussianNB	0.309505	0.467139	0.467139
2	MultinomialNB	0.430356	0.436515	0.436515
3	ComplementNB	0.396336	0.385406	0.385406
4	Decision Tree	0.579231	0.558462	0.542308
5	Random Forest	0.695385	0.660769	0.678462
6	Gradient Boosting	0.634615	0.619231	0.643077

```
In [ ]: #print(dt_raw_set.keys())
#print(dt_raw_set.values(), dt_pca_set.values(), dt_lda_set.values())
dt_s =[dt_raw_set, dt_pca_set, dt_lda_set]
df_dt_settings = pd.DataFrame(dt_s)
df_dt_settings['Algo']= 'Decision Tree'
```

```
df_dt_settings['Data'] = ['raw', 'pca', 'lda']
display(df_dt_settings)
```

In [65]:

```
rf_s =[rf_raw_set,rf_pca_set,rf_lda_set]
df_rf_settings = pd.DataFrame(rf_s)
df_rf_settings['Algo']= 'Random Forest'
df_rf_settings['Data'] = ['raw', 'pca', 'lda']
display(df_rf_settings)
```

	<b>criterion</b>	<b>max_depth</b>	<b>n_estimators</b>	<b>Algo</b>	<b>Data</b>
<b>0</b>	entropy	18	500	Random Forest	raw
<b>1</b>	entropy	16	200	Random Forest	pca
<b>2</b>	entropy	17	500	Random Forest	lda

In [66]:

```
gbc_s =[gbc_set,gbc_pca_set,gbc_lda_set]
df_gbc_settings = pd.DataFrame(gbc_s)
df_gbc_settings['Algo']= 'Gradient Boosting'
df_gbc_settings['Data'] = ['raw', 'pca', 'lda']
display(df_gbc_settings)
```

	<b>learning_rate</b>	<b>n_estimators</b>	<b>Algo</b>	<b>Data</b>
<b>0</b>	0.1	800	Gradient Boosting	raw
<b>1</b>	0.1	800	Gradient Boosting	pca
<b>2</b>	0.1	800	Gradient Boosting	lda

In [ ]:

```
knn_setting = ({'n_neighbors':23 , 'metric': 'manhattan' , 'weights':
"distance"})
gnb_settings = ("default")
mnb_settings = ("default")
cnb_settings = ("default")
dt_settings =
({'criterion':'Entropy','max_depth':18,'min_samples_leaf':2,'min_samples_split'

rf_settings = ({'criterion':'Entropy','max_depth':18,'n_estimators':500})
gb_settings = ({'learning_rate':0.1,'n_estimators':800})

settings_all =
[knn_setting,gnb_settings,mnb_settings,cnb_settings,dt_settings,dt_settings,rf_

settings_all
df_final['settings'] =settings_all
```

```
In [9]: display(df_final)
```

	Algo	settings	wine_raw	wine_pca	wine_lda
0	KNeighborsClassifier	{'n_neighbors': 23, 'metric': 'manhattan', 'we...	0.504236	0.489310	0.489310
1	GaussianNB	default'	0.309505	0.467139	0.467139
2	MultinomialNB	default'	0.430356	0.436515	0.436515
3	ComplementNB	default'	0.396336	0.385406	0.385406
4	Decision Tree	{'criterion': 'Entropy','max_depth': 18,'min_s...	0.579231	0.558462	0.542308
5	Random Forest	{'criterion': 'Entropy', 'max_depth': 18, 'n_e...	0.695385	0.660769	0.678462
6	Gradient Boosting	{'learning_rate': 0.1, 'n_estimators': 800}	0.634615	0.619231	0.643077

## Comments

- Random Forest with following settings *criterion: entropy, max-depth: 18, n-estimators:500* on raw data(MinMax Scaling) as well as for PCA and LDA dataset performed the best task at classification.
- For Gradient Boosting and Gaussian Naive bayes , dimensionality reduction did help while for other algorithms the accuracy results were almost comparable to original dataset

# Abalone Dataset

## Preprocessing

### Loading Data

```
In [28]: #Columns/Features
D = ['Sex', 'Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
     'Viscera weight', 'Shell weight']
L = 'Rings'
DL = D + [L]

#Loading Data set
abalone = pd.read_csv("abalone.csv", sep=',', names=DL)

abalone = abalone.replace('M', 0)
abalone = abalone.replace('F', 1)
abalone = abalone.replace('I', 2)
```

### Normalization/Scaling

```
In [29]: abalone1=abalone.copy()
abalone_z = abalone1.drop(columns = ['Sex', 'Rings'])
abalone_znormalized = abalone_z.apply(stats.zscore)
abalone_znormalized['Sex'] = abalone['Sex']
abalone_znormalized['Rings'] = abalone['Rings']

target_z = abalone_znormalized['Rings']
data_z = abalone_znormalized.drop(columns = 'Rings')
X_train_z, X_test_z, Y_train_z, Y_test_z = train_test_split(data_z, target_z,
test_size=0.2 , random_state=27)
```

```
In [30]: abalone_s = abalone.copy()
abalone_s = abalone_s.drop(columns = ['Sex' , 'Rings'])
abalone_minmax = (abalone_s - abalone_s.min()) / ( abalone_s.max() -
abalone_s.min())
abalone_minmax['Sex'] = abalone['Sex']
abalone_minmax['Rings'] = abalone['Rings']

abalone_min_max = abalone_minmax
target = abalone_min_max['Rings']
data = abalone_min_max.drop(columns = 'Rings')
X_train, X_test, Y_train, Y_test = train_test_split(data, target,
test_size=0.2 , random_state=27)
```

## Best formulated KNN from Assignment-1

```
In [31]: knn = KNeighborsClassifier(n_neighbors=66, metric='minkowski',
weights='uniform', p=2)
knn.fit(X_train, Y_train)
Y_pred = knn.predict(X_test)

knn_train_acc = knn.score(X_train, Y_train)
print("Training Score: ", knn_train_acc)
knn_test_acc = knn.score(X_test, Y_test)
print("Testing Score: ", knn_test_acc)
```

Training Score: 0.2975157138581263

Testing Score: 0.2631578947368421

## Representation Learning

### Indepeendent and dependent variable

```
In [32]: #Scaled 'data' and 'target'
x=data
y=target
n_components=2
target_names = np.sort(y.unique())
```

## PCA

```
In [33]: print("PCA Model - Explained Variance of each component")
pca = PCA(n_components=n_components)
X_r = pca.fit(x).transform(x)
exp_var_ratio_df = pd.DataFrame(pca.explained_variance_ratio_, columns=
['Explained Variance Ratio'])
print(exp_var_ratio_df)

plt.figure(figsize=(7,7))
colors = plt.cm.get_cmap("twilight")


for i, target_name in enumerate(target_names):
    plt.scatter(
        X_r[y == i, 0], X_r[y == i, 1], alpha=0.5, lw=1, label=target_name,
        cmap="twilight"
    )
plt.ylim(-1.5,1.5)
plt.legend(loc="best", shadow=False, scatterpoints=1)
plt.xlabel("Prinicipal Component 1")
plt.ylabel("Prinicipal Component 2")
```

```
plt.title("First Two Components from PCA")
plt.set_cmap("twilight")

plt.show()
```

PCA Model - Explained Variance of each component

	Explained Variance Ratio
0	0.865084
1	0.123693



## LDA

In [34]:

```
print("LDA Model - Explained Variance of each component")
lda = LinearDiscriminantAnalysis(n_components=n_components)
X_r2 = lda.fit(x, y).transform(x)
exp_var_ratio_df = pd.DataFrame(lda.explained_variance_ratio_, columns=
['Explained Variance Ratio'])
print(exp_var_ratio_df)


plt.figure(figsize=(7,7))
colors = plt.cm.get_cmap("twilight")
for i, target_name in enumerate(target_names):
    plt.scatter(
        X_r2[y == i, 0], X_r2[y == i, 1], alpha=0.5, lw=1, label=target_name,
        cmap="twilight"
    )
```

```
#plt.xlim(-4,4)
plt.legend(loc="best", shadow=False, scatterpoints=1)
plt.xlabel("Component 1")
plt.ylabel("Component 2")
plt.title("First Two Components from LDA")
plt.set_cmap("twilight")

plt.show()
```

LDA Model - Explained Variance of each component

	Explained Variance Ratio
0	0.719139
1	0.216991



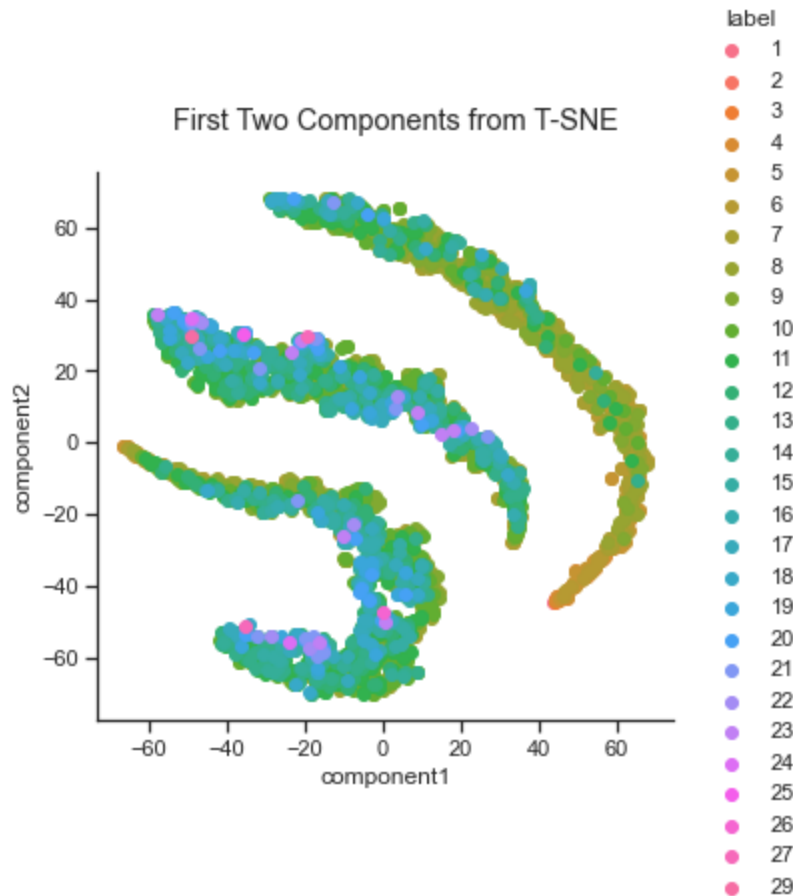
## T-SNE

In [35]:

```
tsne = TSNE(n_components=n_components)
X_r3 = tsne.fit_transform(x)
tsne_data = np.vstack((X_r3.T, target)).T
tsne_df = pd.DataFrame(data = tsne_data, columns = ['component1',
'component2', 'label'] )
#print(tsne_df.head())
tsne_df['label'] = tsne_df['label'].astype('int')
tsne_df['label'] = tsne_df['label'].astype('category')
g = sns.FacetGrid(tsne_df, hue='label', height=5).map(plt.scatter,
'component1', 'component2')
```



```
g.add_legend()
g.fig.suptitle("First Two Components from T-SNE ")
g.tight_layout()
```



## Comments

- PCA: Three distant lines due to presence of categorical column(Sex) in the data. We have one principal component that explains nearly all (85%) of the variance the reason is, the principal component is a mixture of all variables in nearly equal proportions. That is, the principal component is the one that combines every single variable.
- LDA: Inter Class spread is increased and intra-class spread is decreased, clustering similar class points together. Outliers evident in the plot as distance between them is large.
- t-SNE: Keeps changing with the run, three clusters are formed, better way to visualize but does not help in our case for multiclass classification

## Scree Plot (PCA)

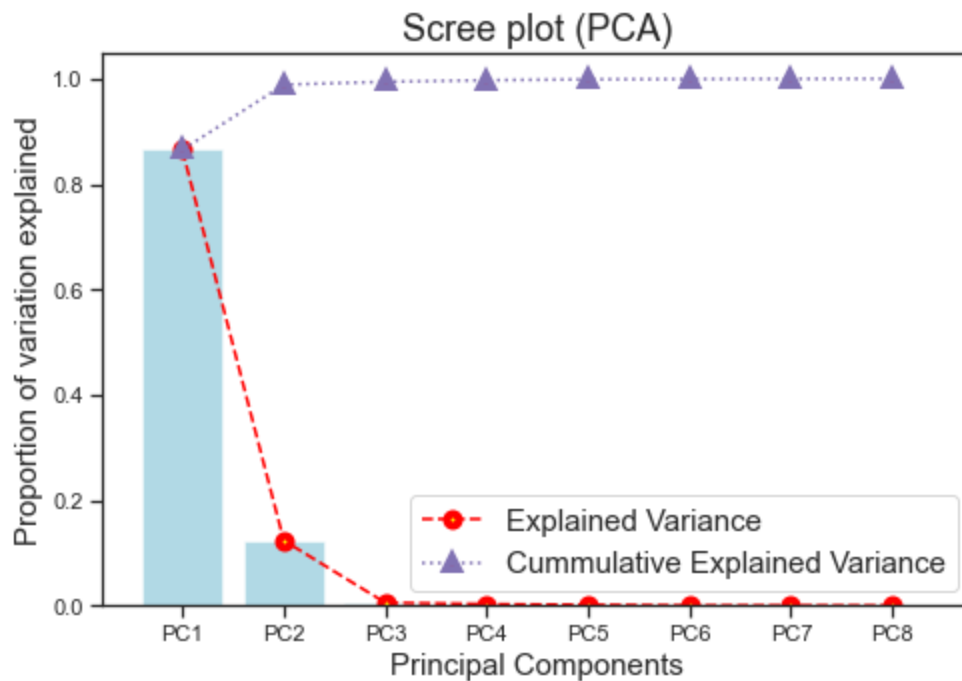
In [36]:

```
pca = PCA(n_components=None)
X_r = pca.fit(x).transform(x)
explained_variance_pca = pca.explained_variance_ratio_
fig = plt.figure(figsize=(7,5))
tick_label = ['PC' + str(i) for i in range(1,len(explained_variance_pca)+1)]
plt.bar(range(len(explained_variance_pca)), explained_variance_pca, color =
'c',alpha=0.5, align='center',tick_label=tick_label)
plt.plot(range(len(explained_variance_pca)), explained_variance_pca, color =
'red',marker='o',mew=4,mfc='yellow',ls='--')
```

```

,linewidth=1.5,label='Explained Variance')
plt.plot(range(len(explained_variance_pca)), explained_variance_pca.cumsum(),
color = 'm',marker = '^',mew =4,mfc='red',ls = ':')
,linewidth=1.5,label='Cummulative Explained Variance')
plt.ylabel('Proportion of variation explained',fontsize = 15)
plt.xlabel('Principal Components',fontsize = 15)
plt.title('Scree plot (PCA)',fontsize = 18)
plt.legend(loc='best',fontsize = 15)
#plt.axis("off")
plt.tight_layout()
plt.show()

```



## Reduced Dimensionality

### PCA

In [37]:

```

knn = KNeighborsClassifier(n_neighbors=66, metric='minkowski',
weights='uniform', p=2)
D = data.shape[1]
#print(D)
acc_list =[]

for i in range(2,D):
    pca = PCA(n_components=i)
    X_train_r = pca.fit(X_train).transform(X_train)
    X_test_r = pca.transform(X_test)
    knn.fit(X_train_r, Y_train)
    y_pred = knn.predict(X_test_r)
    acc_list.append(metrics.accuracy_score(Y_test,y_pred))


```

```

fig = plt.figure(figsize=(7,5))
plt.plot(range(2,D),acc_list, color = 'blue',marker = 'o',mew =4,mfc='red',ls
='--' ,linewidth=1.5,label='Principal Components')
plt.ylabel('Accuracy Score',fontsize = 15)
plt.xlabel('Principal Components',fontsize = 15)
plt.title('Accuracy vs # Principal Components',fontsize = 18)
plt.legend(loc='best',fontsize = 15)
plt.tight_layout()
plt.show()

print("The maximum accuracy {} happens to be when number of Principal
Components is : {}".format(max(acc_list),acc_list.index(max(acc_list))+2))

```



The maximum accuracy 0.27392344497607657 happens to be when number of Principal Components is : 7

## LDA

In [38]:

```

knn = KNeighborsClassifier(n_neighbors=66, metric='minkowski',
weights='uniform', p=2)
#LDA is constrained by the number of components it can use , given by :
min(n_classes-1, len(features))
D = min(data.shape[1],len(target_names))
acc_list =[]

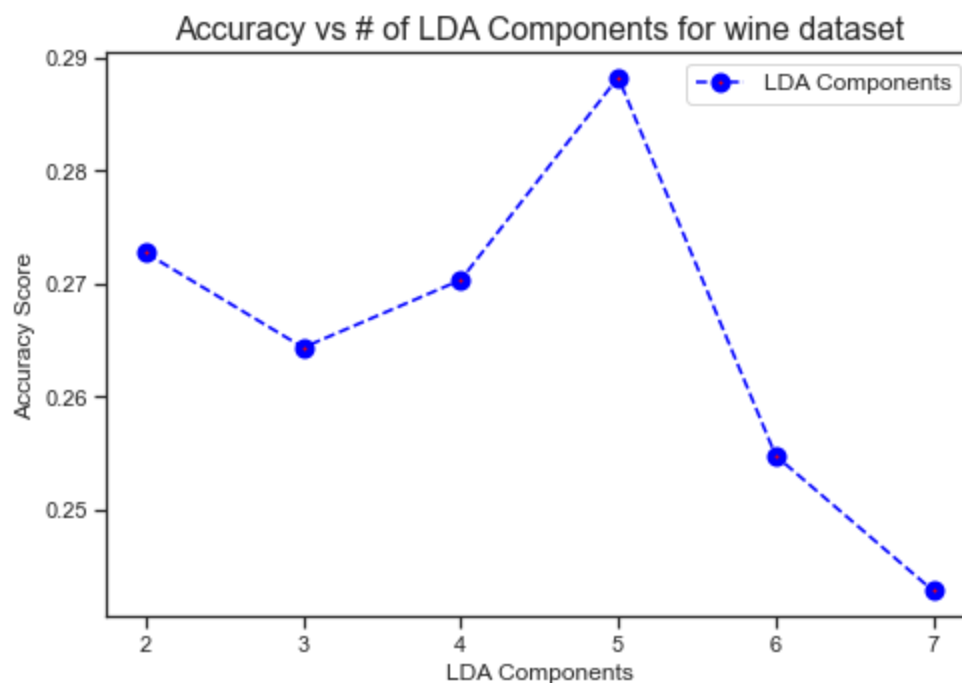
for i in range(2,D):
    lda = LinearDiscriminantAnalysis(n_components=i)
    X_train_r = lda.fit(X_train,Y_train).transform(X_train)
    X_test_r = lda.transform(X_test)
    knn.fit(X_train_r, Y_train)
    y_pred = knn.predict(X_test_r)

```

```
acc_list.append(metrics.accuracy_score(Y_test,y_pred))
```

```
fig = plt.figure(figsize=(7,5))
plt.plot(range(2,D),acc_list, color = 'blue',marker = 'o',mew =4,mfc='red',ls
='--' ,linewidth=1.5,label='LDA Components')
plt.ylabel('Accuracy Score',fontsize = 12)
plt.xlabel('LDA Components',fontsize = 12)
plt.title('Accuracy vs # of LDA Components for wine dataset',fontsize = 16)
plt.legend(loc='best',fontsize = 12)
plt.tight_layout()
plt.show()

print("The maximum accuracy {} happens to be when number of LDA Components is
: {}".format(max(acc_list),acc_list.index(max(acc_list))+2))
```



The maximum accuracy 0.28827751196172247 happens to be when number of LDA Components is : 5

## Comments

- Two Principal Components are able to explain 98% of the variance in the dataset. We get maximum accuracy with 7 principal components which is same as original dimension of dataset, but we use 5 components in subsequent sections as there is no significant difference in accuracy with 5 and 7 components.
- With 5 components in LDA, we get maximum accuracy (more than PCA). LDA does a better job at classification task.
- Dimensionality Reduction improves the accuracy for both PCA and LDA compared to best formulation from Assignment-1

## Versions of dataset (raw/pca/lda)

```
In [39]: abalone_raw = abalone.copy()
```

```

abalone_normalized_raw = abalone_znormalized.copy()
abalone_minmax_normalized = abalone_minmax.copy()

pca = PCA(n_components=5)
X_r = pca.fit(x).transform(x)
pca_data = np.vstack((X_r.T, target)).T
abalone_pca = pd.DataFrame(data = pca_data, columns =
['component1', 'component2', 'component3', 'component4', 'component5', 'label'] )
abalone_temp_pca = abalone_pca.copy()
abalone_temp_pca_2 = abalone_temp_pca.drop(columns = ['label'])
abalone_pca = (abalone_temp_pca_2 - abalone_temp_pca_2.min()) / (
abalone_temp_pca_2.max() - abalone_temp_pca_2.min())
abalone_pca['label'] = abalone_temp_pca['label']
#abalone_pca = (abalone_pca - abalone_pca.min()) / ( abalone_pca.max() -
abalone_pca.min())
abalone_pca['label'] = abalone_pca['label'].astype('int')
abalone_pca['label'] = abalone_pca['label'].astype('category')
#print(abalone_pca.head())

lda = LinearDiscriminantAnalysis(n_components=5)
X_r2 = lda.fit(x, y).transform(x)
lda_data = np.vstack((X_r2.T, target)).T
abalone_lda = pd.DataFrame(data = lda_data, columns = ['component1',
'component2', 'component3', 'component4', 'component5', 'label'] )
abalone_temp_lda = abalone_lda.copy()
abalone_temp_lda_2 = abalone_temp_lda.drop(columns = ['label'])
abalone_lda = (abalone_temp_lda_2 - abalone_temp_lda_2.min()) / (
abalone_temp_lda_2.max() - abalone_temp_lda_2.min())
abalone_lda['label'] = abalone_temp_lda['label']
abalone_lda['label'] = abalone_lda['label'].astype('int')
abalone_lda['label'] = abalone_lda['label'].astype('category')
#print(abalone_lda.head())

```

In [40]:

```

abalone_p = abalone_pca.copy()
target_pca = abalone_p['label']
data_pca = abalone_p.drop(columns = 'label')
X_train_pca, X_test_pca, Y_train_pca, Y_test_pca = train_test_split(data_pca,
target_pca, test_size=0.2 , random_state=27)

abalone_l = abalone_lda.copy()
target_lda = abalone_l['label']
data_lda = abalone_l.drop(columns = 'label')

```

```
X_train_lda, X_test_lda, Y_train_lda, Y_test_lda = train_test_split(data_lda,
target_lda, test_size=0.2 , random_state=27)
```

## Naive Bayes Classifier

### 5-Fold CV for Naive Bayes & KNN classifier

In [64]:

```
kfold = KFold(n_splits=5)
clrs = []

clrs.append(KNeighborsClassifier(n_neighbors=23 , metric = 'manhattan' ,
weights = "distance"))
clrs.append(GaussianNB())
clrs.append(MultinomialNB())
clrs.append(ComplementNB())

cv_results = []
cv_results_pca = []
cv_results_lda = []

for clr in clrs :
    cv_results.append(cross_val_score(clr, data, target , scoring =
'accuracy', cv = kfold, n_jobs=-1))
    cv_results_pca.append(cross_val_score(clr, data_pca, target_pca , scoring
= 'accuracy', cv = kfold, n_jobs=-1))
    cv_results_lda.append(cross_val_score(clr, data_lda, target_lda , scoring
= 'accuracy', cv = kfold, n_jobs=-1))

cv_means = []
cv_means_pca = []
cv_means_lda = []

for cv_result in cv_results:
    cv_means.append(cv_result.mean())

for cv_result in cv_results_pca:
    cv_means_pca.append(cv_result.mean())

for cv_result in cv_results_lda:
    cv_means_lda.append(cv_result.mean())

cv_df = pd.DataFrame({"Mean_Accuracy":cv_means,"Algo":
['KNeighborsClassifier','GaussianNB','MultinomialNB','ComplementNB']})
cv_df_pca = pd.DataFrame({"Mean_Accuracy":cv_means_pca,"Algo":
```

```
[ 'KNeighborsClassifier', 'GaussianNB', 'MultinomialNB', 'ComplementNB' ]})
cv_df_lda = pd.DataFrame({"Mean_Accuracy":cv_means_lda,"Algo":
[ 'KNeighborsClassifier', 'GaussianNB', 'MultinomialNB', 'ComplementNB' ]})

#print(cv_df)
#print(cv_df_pca)
#print(cv_df_lda)
```

In [65]:

```
data_frames = [cv_df, cv_df_pca, cv_df_lda]
df_merged = reduce(lambda left,right: pd.merge(left,right,on=
['Algo'],how='outer'), data_frames)
df_merged.rename(columns={"Mean_Accuracy_x": "abalone_raw",
"Mean_Accuracy_y": "abalone_pca","Mean_Accuracy":"abalone_lda"},inplace=True)
df_merged.set_index('Algo',inplace=True)
print(df_merged)
```

	abalone_raw	abalone_pca	abalone_lda
Algo			
KNeighborsClassifier	0.250907	0.249947	0.253777
GaussianNB	0.230326	0.232716	0.251631
MultinomialNB	0.206377	0.165673	0.164955
ComplementNB	0.170705	0.171898	0.228406

## Comments

- Complement Naive Bayes performs poorly, the dataset is unbalanced but due to multiclass classification it does not perform as expected for unbalanced datasets.

## Decision Trees Classifier

### 5-Fold CV and hyperparameter tuning

In [45]:

```
dtc = DecisionTreeClassifier()
dtc.fit(X_train, Y_train)
Y_pred = dtc.predict(X_test)
#cm=confusion_matrix(Y_test, Y_pred)
#print(confusion_matrix(Y_test, Y_pred))
#print(classification_report(Y_test, Y_pred))
print("Accuracy Score for Decision Tree classifier without hyper parameter
tuning is: ",accuracy_score(Y_test, Y_pred))
#disp =
ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=dtc.classes_)
#disp.plot()

grid_params = {
    #'criterion' : ['gini', 'entropy'],
    'max_depth' : range(2,25,1),
```

```

'min_samples_split' : range(2, 5, 1),
'min_samples_leaf' : range(2, 5, 1)
}

grid_search = GridSearchCV(dtc, grid_params, cv = 5, n_jobs = -1, verbose =
1)
grid_search.fit(X_train, Y_train)
dtc = grid_search.best_estimator_
Y_pred = dtc.predict(X_test)
print("Accuracy Score is :",accuracy_score(Y_test, Y_pred))
# best parameters and best score
print("Best Parameters are : ",grid_search.best_params_)
#print("Best Score is : ",grid_search.best_score_)
dt_raw_acc = accuracy_score(Y_test, Y_pred)
dt_raw_set=grid_search.best_params_

grid_search.fit(X_train_pca, Y_train_pca)
dtc = grid_search.best_estimator_
Y_pred_pca = dtc.predict(X_test_pca)
print("Accuracy Score using PCA is :",accuracy_score(Y_test_pca, Y_pred_pca))
# best parameters and best score
print("Best Parameters using PCA are : ",grid_search.best_params_)
#print("Best Score using PCA is : ",grid_search.best_score_)
dt_pca_acc = accuracy_score(Y_test_pca, Y_pred_pca)
dt_pca_set=grid_search.best_params_

grid_search.fit(X_train_lda, Y_train_lda)
dtc = grid_search.best_estimator_
Y_pred_lda = dtc.predict(X_test_lda)
print("Accuracy Score using LDA is :",accuracy_score(Y_test_lda, Y_pred_lda))
# best parameters and best score
print("Best Parameters using LDA are : ",grid_search.best_params_)
#print("Best Score using LDA is : ",grid_search.best_score_)
dt_lda_acc = accuracy_score(Y_test_lda, Y_pred_lda)
dt_lda_set=grid_search.best_params_

```

Accuracy Score for Decision Tree classifier without hyper parameter tuning is: 0.19617224880382775

Fitting 5 folds for each of 207 candidates, totalling 1035 fits

Accuracy Score is : 0.2535885167464115

Best Parameters are : {'max\_depth': 4, 'min\_samples\_leaf': 2, 'min\_samples\_split': 2}

Fitting 5 folds for each of 207 candidates, totalling 1035 fits

Accuracy Score using PCA is : 0.26435406698564595

Best Parameters using PCA are : {'max\_depth': 4, 'min\_samples\_leaf': 4, 'min\_samples\_split': 2}

Fitting 5 folds for each of 207 candidates, totalling 1035 fits

Accuracy Score using LDA is : 0.23444976076555024



```
Best Parameters using LDA are : {'max_depth': 3, 'min_samples_leaf': 2, 'min_samples_split': 2}
```

```
In [46]: dt_raw_acc = accuracy_score(Y_test, Y_pred)
dt_pca_acc = accuracy_score(Y_test_pca, Y_pred_pca)
dt_lda_acc = accuracy_score(Y_test_lda, Y_pred_lda)
# print(dt_raw_acc, dt_pca_acc, dt_lda_acc)
```


```
In [47]: cm=confusion_matrix(Y_test, Y_pred)
cm_pca=confusion_matrix(Y_test_pca, Y_pred_pca)
cm_lda=confusion_matrix(Y_test_lda, Y_pred_lda)
cm_dict = {
    "raw": cm,
    "pca": cm_pca,
    "lda": cm_lda
}

f, axes = plt.subplots(1, 3, figsize=(20, 5), sharey='row')
for i, (c, cm) in enumerate(cm_dict.items()):
    #print(i, cm)
    disp = ConfusionMatrixDisplay(cm)
    disp.plot(ax=axes[i], xticks_rotation=45)
    disp.ax_.set_title(c)
    disp.im_.colorbar.remove()
    disp.ax_.set_xlabel('')
    if i!=0:
        disp.ax_.set_ylabel('')

f.text(0.4, 0.1, 'Predicted label', ha='left')
plt.subplots_adjust(wspace=0.40, hspace=0.1)

f.suptitle('Confusion matrix', fontsize=12)
f.colorbar(disp.im_, ax=axes)
plt.show()
```

### Confusion matrix



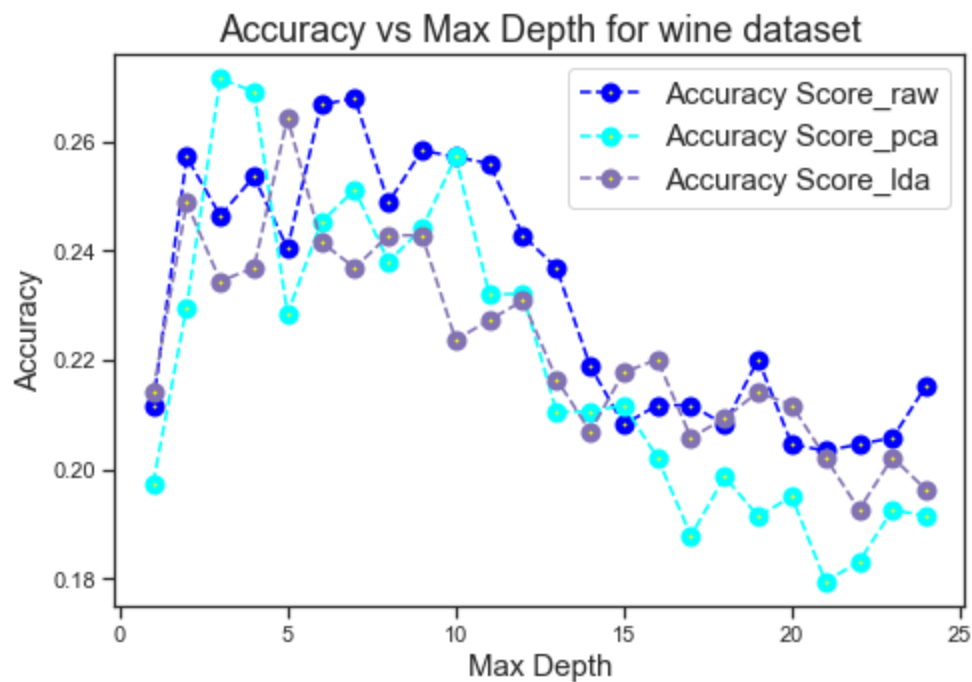
## Accuracy vs Max\_depth

In [48]:

```
def plot_acc_vs_depth(dataset):
    max_depth_range = list(range(1, 25))
    acc_list = []
    acc_list_pca = []
    acc_list_lda = []
    for depth in max_depth_range:
        dtc_w = DecisionTreeClassifier(max_depth = depth)
        if dataset[0] == 'raw':
            dtc_w.fit(X_train, Y_train)
            score = dtc_w.score(X_test, Y_test)
            acc_list.append(score)
        if dataset[1] == 'pca':
            dtc_w.fit(X_train_pca, Y_train_pca)
            score = dtc_w.score(X_test_pca, Y_test_pca)
            acc_list_pca.append(score)
        if dataset[2] == 'lda':
            dtc_w.fit(X_train_lda, Y_train_lda)
            score = dtc_w.score(X_test_lda, Y_test_lda)
            acc_list_lda.append(score)

    return (acc_list, acc_list_pca, acc_list_lda)


acc_list, acc_list_pca, acc_list_lda = plot_acc_vs_depth(['raw', 'pca', 'lda'])
fig = plt.figure(figsize=(7,5))
plt.plot(range(1,25),acc_list, color = 'blue',marker = 'o',mew
=4,mfc='yellow',ls = '--' ,linewidth=1.5,label='Accuracy Score_raw')
plt.plot(range(1,25),acc_list_pca, color = 'cyan',marker = 'o',mew
=4,mfc='yellow',ls = '--' ,linewidth=1.5,label='Accuracy Score_pca')
plt.plot(range(1,25),acc_list_lda, color = 'm',marker = 'o',mew
=4,mfc='yellow',ls = '--' ,linewidth=1.5,label='Accuracy Score_lda')
plt.ylabel('Accuracy',fontsize = 15)
plt.xlabel('Max Depth',fontsize = 15)
plt.title('Accuracy vs Max Depth for wine dataset',fontsize = 18)
plt.legend(loc='best',fontsize = 15)
plt.tight_layout()
plt.show()
```



## Splitting rules used for the trees

```
In [49]: dtc_g = DecisionTreeClassifier(criterion= 'entropy', max_depth=4,
min_samples_leaf= 2, min_samples_split=2)
dtc_g.fit(X_train,Y_train)
fn=X_train.columns
cn = abalone_raw.Rings.unique().astype(str)
dot_data = tree.export_graphviz(dtc_g,
filled=True,feature_names=fn,class_names=cn ,rounded=True,
special_characters=True)
graph = graphviz.Source(dot_data)
graph
```

Out[49]:



```
In [60]: from sklearn.tree import export_text
dtc_g = DecisionTreeClassifier(criterion= 'entropy', max_depth=4,
min_samples_leaf= 2, min_samples_split=2)
dt_rules = dtc_g.fit(X_train,Y_train)
fn_t=X_train.columns
#print(type(fn_t.tolist()))
tree_rules = export_text(dt_rules,feature_names = fn_t.tolist())
print(tree_rules)
```

```

|--- Shell weight <= 0.15
|   |--- Shell weight <= 0.05
|   |   |--- Shell weight <= 0.02
|   |   |   |--- Shucked weight <= 0.00

```

```

| | | | |---- class: 3
| | | | |---- Shucked weight > 0.00
| | | | |---- class: 4
| | | |---- Shell weight > 0.02
| | | | |---- Shucked weight <= 0.04
| | | | |---- class: 5
| | | | |---- Shucked weight > 0.04
| | | | |---- class: 6
| | |---- Shell weight > 0.05
| | |---- Sex <= 1.50
| | | | |---- Shucked weight <= 0.16
| | | | |---- class: 9
| | | | |---- Shucked weight > 0.16
| | | | |---- class: 8
| | |---- Sex > 1.50
| | | | |---- Shell weight <= 0.12
| | | | |---- class: 7
| | | | |---- Shell weight > 0.12
| | | | |---- class: 8
|---- Shell weight > 0.15
| |---- Shell weight <= 0.29
| | |---- Shucked weight <= 0.26
| | | | |---- Shell weight <= 0.19
| | | | |---- class: 8
| | | | |---- Shell weight > 0.19
| | | | |---- class: 10
| | | |---- Shucked weight > 0.26
| | | | |---- Shell weight <= 0.26
| | | | |---- class: 9
| | | | |---- Shell weight > 0.26
| | | | |---- class: 9
| |---- Shell weight > 0.29
| | |---- Shucked weight <= 0.30
| | | | |---- Height <= 0.17
| | | | |---- class: 10
| | | | |---- Height > 0.17
| | | | |---- class: 13
| | |---- Shucked weight > 0.30
| | | | |---- Shell weight <= 0.41
| | | | |---- class: 10
| | | | |---- Shell weight > 0.41
| | | | |---- class: 11

```

## Comments

- As observed in Assignment-1, weight columns were positively correlated with many other columns (namely; Height, Length, Diameter) and hence it is chosen as the root node for the split. Weight and Sex are the major contributors to splitting rules indicating that these are more important for classification.

## Splitting rules used for the trees but fitted on page

```

In [51]: dtc_g = DecisionTreeClassifier(criterion= 'entropy', max_depth=4,
min_samples_leaf= 2, min_samples_split=5)
dtc_g.fit(X_train,Y_train)
fn=X_train.columns

```

```

cn = abalone_raw.Rings.unique().astype(str)
fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (8,8), dpi=300)
tree.plot_tree(dtc_g, filled=True,feature_names=fn,class_names=cn
,rounded=True)
#fig.savefig('dtc_tree_plot_3.png')

```

Out[51]:


```

[Text(930.0, 1630.8, 'Shell weight <= 0.15\nentropy = 3.593\nsamples = 3341\nvalue = [1,
1, 10, 46, 87, 202, 310, 453, 557, 520, 387\n212, 166, 98, 77, 54, 50, 33, 27, 22, 13, 3,
5\n2, 1, 1, 2, 1]\nnclass = 14'),
Text(465.0, 1268.4, 'Shell weight <= 0.046\nentropy = 2.972\nsamples = 980\nvalue = [1,
1, 10, 46, 87, 189, 249, 174, 104, 62, 19, 18\n11, 3, 2, 3, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0]\nnclass = 16'),
Text(232.5, 906.0, 'Shell weight <= 0.02\nentropy = 2.386\nsamples = 219\nvalue = [1, 1,
10, 45, 67, 59, 26, 8, 1, 1, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0]\nnclass
= 8'),
Text(116.25, 543.6000000000001, 'Diameter <= 0.08\nentropy = 2.028\nsamples = 74\nvalue =
[1, 1, 10, 35, 18, 7, 2, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nnc
lass = 10'),
Text(58.125, 181.20000000000005, 'entropy = 1.5\nsamples = 4\nvalue = [1, 1, 2, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nnclass = 9'),
Text(174.375, 181.20000000000005, 'entropy = 1.84\nsamples = 70\nvalue = [0, 0, 8, 35, 1
8, 7, 2, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nnclass = 10'),
Text(348.75, 543.6000000000001, 'Shucked weight <= 0.036\nentropy = 2.085\nsamples = 145
\nvalue = [0, 0, 0, 10, 49, 52, 24, 8, 1, 1, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0]\nnclass = 20'),
Text(290.625, 181.20000000000005, 'entropy = 2.179\nsamples = 94\nvalue = [0, 0, 0, 9, 3
6, 28, 11, 8, 1, 1, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nnclass = 8'),
Text(406.875, 181.20000000000005, 'entropy = 1.628\nsamples = 51\nvalue = [0, 0, 0, 1, 1
3, 24, 13, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nnclass = 20'),
Text(697.5, 906.0, 'Sex <= 1.5\nentropy = 2.714\nsamples = 761\nvalue = [0, 0, 0, 1, 20,
130, 223, 166, 103, 61, 19, 18\n11, 3, 2, 3, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0]\nnclass =
16'),
Text(581.25, 543.6000000000001, 'Shucked weight <= 0.156\nentropy = 2.905\nsamples = 280
\nvalue = [0, 0, 0, 0, 3, 21, 48, 58, 62, 46, 15, 13, 8\n2, 2, 1, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0]\nnclass = 14'),
Text(523.125, 181.20000000000005, 'entropy = 2.952\nsamples = 234\nvalue = [0, 0, 0, 0,
3, 15, 34, 38, 59, 43, 15, 13, 8\n2, 2, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0]\nnclass = 1
4'),
Text(639.375, 181.20000000000005, 'entropy = 1.942\nsamples = 46\nvalue = [0, 0, 0, 0, 0,
6, 14, 20, 3, 3, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nnclass = 19'),
Text(813.75, 543.6000000000001, 'Shell weight <= 0.118\nentropy = 2.371\nsamples = 481\nv
alue = [0, 0, 0, 1, 17, 109, 175, 108, 41, 15, 4, 5, 3\n1, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0]\nnclass = 16'),
Text(755.625, 181.20000000000005, 'entropy = 2.266\nsamples = 355\nvalue = [0, 0, 0, 1, 1
7, 101, 137, 61, 19, 10, 4, 3, 1\n1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0]\nnclass = 1
6'),
Text(871.875, 181.20000000000005, 'entropy = 2.214\nsamples = 126\nvalue = [0, 0, 0, 0,
0, 8, 38, 47, 22, 5, 0, 2, 2, 0\n0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nnclass = 19'),
Text(1395.0, 1268.4, 'Shell weight <= 0.29\nentropy = 3.336\nsamples = 2361\nvalue = [0,
0, 0, 0, 13, 61, 279, 453, 458, 368, 194\n155, 95, 75, 51, 50, 32, 27, 22, 13, 3, 5, 2
\n1, 1, 2, 1]\nnclass = 11'),
Text(1162.5, 906.0, 'Shucked weight <= 0.256\nentropy = 3.083\nsamples = 1206\nvalue =
[0, 0, 0, 0, 0, 10, 58, 246, 300, 212, 129, 73\n52, 40, 31, 15, 14, 8, 8, 3, 3, 2, 2, 0, 0
\n0, 0, 0]\nnclass = 14'),
Text(1046.25, 543.6000000000001, 'Shell weight <= 0.192\nentropy = 3.276\nsamples = 801\n
value = [0, 0, 0, 0, 0, 9, 39, 153, 159, 135, 85, 61\n49, 36, 28, 13, 11, 7, 7, 3, 2, 2,
2, 0, 0\n0, 0, 0]\nnclass = 14'),

```

```

Text(988.125, 181.20000000000005, 'entropy = 2.912\nsamples = 325\nvalue = [0, 0, 0, 0,
0, 5, 35, 93, 73, 44, 30, 13, 8\n8, 8, 3, 1, 2, 1, 0, 1, 0, 0, 0, 0, 0\n0]\n\nclass = 1
9'),
Text(1104.375, 181.20000000000005, 'entropy = 3.339\nsamples = 476\nvalue = [0, 0, 0, 0,
0, 4, 4, 60, 86, 91, 55, 48, 41\n28, 20, 10, 10, 5, 6, 3, 1, 2, 2, 0, 0, 0, 0\n0]\n\nclass = 1
11'),
Text(1278.75, 543.6000000000001, 'Shell weight <= 0.257\nentropy = 2.525\nsamples = 405\n
value = [0, 0, 0, 0, 0, 1, 19, 93, 141, 77, 44, 12, 3\n4, 3, 2, 3, 1, 1, 0, 1, 0, 0, 0, 0,
0, 0\n0]\n\nclass = 14'),
Text(1220.625, 181.20000000000005, 'entropy = 2.104\nsamples = 207\nvalue = [0, 0, 0, 0,
0, 0, 15, 65, 80, 32, 10, 4, 0\n0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0]\n\nclass = 1
4'),
Text(1336.875, 181.20000000000005, 'entropy = 2.73\nsamples = 198\nvalue = [0, 0, 0, 0,
0, 1, 4, 28, 61, 45, 34, 8, 3, 4\n3, 2, 2, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0]\n\nclass = 14'),
Text(1627.5, 906.0, 'Shucked weight <= 0.298\nentropy = 3.335\nsamples = 1155\nvalue =
[0, 0, 0, 0, 0, 3, 3, 33, 153, 246, 239, 121\n103, 55, 44, 36, 36, 24, 19, 19, 10, 1, 3, 2
\n1, 1, 2, 1]\n\nclass = 11'),
Text(1511.25, 543.6000000000001, 'Height <= 0.166\nentropy = 3.758\nsamples = 227\nvalue
= [0, 0, 0, 0, 0, 2, 0, 5, 15, 31, 19, 19, 31\n22, 20, 15, 14, 7, 10, 8, 3, 0, 3, 0, 1, 1
\n1, 0]\n\nclass = 11'),
Text(1453.125, 181.20000000000005, 'entropy = 3.672\nsamples = 202\nvalue = [0, 0, 0, 0,
0, 2, 0, 4, 15, 31, 19, 19, 27\n18, 17, 11, 10, 6, 10, 8, 3, 0, 1, 0, 0, 0\n1, 0]\n\nclass = 1
11'),
Text(1569.375, 181.20000000000005, 'entropy = 3.094\nsamples = 25\nvalue = [0, 0, 0, 0,
0, 0, 0, 1, 0, 0, 0, 0, 4, 4\n3, 4, 4, 1, 0, 0, 0, 0, 2, 0, 1, 1, 0, 0]\n\nclass = 13'),
Text(1743.75, 543.6000000000001, 'Shell weight <= 0.407\nentropy = 3.126\nsamples = 928\n
value = [0, 0, 0, 0, 0, 1, 3, 28, 138, 215, 220, 102\n72, 33, 24, 21, 22, 17, 9, 11, 7, 1,
0, 2, 0\n0, 1, 1]\n\nclass = 12'),
Text(1685.625, 181.20000000000005, 'entropy = 2.702\nsamples = 571\nvalue = [0, 0, 0, 0,
0, 1, 2, 27, 120, 165, 133, 51, 35\n15, 8, 1, 4, 5, 1, 1, 2, 0, 0, 0, 0, 0\n0]\n\nclass = 1
11'),
Text(1801.875, 181.20000000000005, 'entropy = 3.398\nsamples = 357\nvalue = [0, 0, 0, 0,
0, 0, 1, 1, 18, 50, 87, 51, 37\n18, 16, 20, 18, 12, 8, 10, 5, 1, 0, 2, 0, 0\n1, 1]\n\nclass
= 12')]
```



## Random Forest Classifier

### 5-Fold CV and hyperparameter tuning

In [52]:

```
rfc = RandomForestClassifier()
rfc.fit(X_train, Y_train)
Y_pred = rfc.predict(X_test)
cm=confusion_matrix(Y_test, Y_pred)
#print(confusion_matrix(Y_test, Y_pred))
#print(classification_report(Y_test, Y_pred))
print("Accuracy Score for Forest classifier without hyper parameter tuning
is: ",accuracy_score(Y_test, Y_pred))
#disp =
ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=dtc.classes_)
#disp.plot()
```

```

grid_params = {
    'n_estimators': [2,5,200, 500],
    'criterion' : ['entropy'],
    'max_depth' : range(2,25,1),
    #'max_features': range(3,10,1),
    #'min_samples_split' : range(2, 10, 1),
    #'min_samples_leaf' : range(2, 10, 1)
}

grid_search = GridSearchCV(rfc, grid_params, cv = 5, n_jobs = -1, verbose =
True)
#print(grid_search.get_params().keys())
grid_search.fit(X_train, Y_train)
rfc = grid_search.best_estimator_
Y_pred = rfc.predict(X_test)
print("Accuracy Score is :",accuracy_score(Y_test, Y_pred))
rf_raw_acc=accuracy_score(Y_test, Y_pred)
# best parameters and best score
print("Best Parameters are : ",grid_search.best_params_)
rf_raw_set=grid_search.best_params_
#print("Best Score is : ",grid_search.best_score_)
ac_df=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
ac_df["accuracy"]=grid_search.cv_results_['mean_test_score']

grid_search.fit(X_train_pca, Y_train_pca)
rfc_pca = grid_search.best_estimator_
Y_pred_pca = rfc_pca.predict(X_test_pca)
print("Accuracy Score using PCA is :",accuracy_score(Y_test_pca, Y_pred_pca))
rf_pca_acc=accuracy_score(Y_test_pca, Y_pred_pca)
# best parameters and best score
print("Best Parameters using PCA are : ",grid_search.best_params_)
rf_pca_set=grid_search.best_params_
#print("Best Score using PCA is : ",grid_search.best_score_)
ac_df_pca=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
ac_df_pca["accuracy"]=grid_search.cv_results_['mean_test_score']

grid_search.fit(X_train_lda, Y_train_lda)
rfc_lda = grid_search.best_estimator_
Y_pred_lda = rfc_lda.predict(X_test_lda)
print("Accuracy Score using LDA is :",accuracy_score(Y_test_lda, Y_pred_lda))
rf_lda_acc=accuracy_score(Y_test_lda, Y_pred_lda)

```



```

# best parameters and best score
print("Best Parameters using LDA are : ",grid_search.best_params_)
#print("Best Score using LDA is : ",grid_search.best_score_)
rf_lda_set=grid_search.best_params_
ac_df_lda=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
ac_df_lda["accuracy"]=grid_search.cv_results_['mean_test_score']

```

Accuracy Score for Forest classifier without hyper parameter tuning is: 0.24641148325358853

Fitting 5 folds for each of 92 candidates, totalling 460 fits

Accuracy Score is : 0.27392344497607657

Best Parameters are : {'criterion': 'entropy', 'max\_depth': 5, 'n\_estimators': 5}

Fitting 5 folds for each of 92 candidates, totalling 460 fits

Accuracy Score using PCA is : 0.2679425837320574

Best Parameters using PCA are : {'criterion': 'entropy', 'max\_depth': 6, 'n\_estimators': 200}

Fitting 5 folds for each of 92 candidates, totalling 460 fits

Accuracy Score using LDA is : 0.2727272727272727

Best Parameters using LDA are : {'criterion': 'entropy', 'max\_depth': 4, 'n\_estimators': 500}

In [53]:

```


cm=confusion_matrix(Y_test, Y_pred)
cm_pca=confusion_matrix(Y_test_pca, Y_pred_pca)
cm_lda=confusion_matrix(Y_test_lda, Y_pred_lda)
cm_dict = {
    "raw": cm,
    "pca": cm_pca,
    "lda": cm_lda
}

f, axes = plt.subplots(1, 3, figsize=(20, 5), sharey='row')
for i, (c, cm) in enumerate(cm_dict.items()):
    #print(i, cm)
    disp = ConfusionMatrixDisplay(cm)
    disp.plot(ax=axes[i], xticks_rotation=45)
    disp.ax_.set_title(c)
    disp.im_.colorbar.remove()
    disp.ax_.set_xlabel('')
    if i!=0:
        disp.ax_.set_ylabel('')

f.text(0.4, 0.1, 'Predicted label', ha='left')
plt.subplots_adjust(wspace=0.40, hspace=0.1)

```

```
f.colorbar(dispatch.im_, ax=axes)
plt.show()
```



## Heatmap of max\_depth vs n\_estimators


In [54]:

```
#Pivoting the dataframe for plotting heat map
ac_df=ac_df.pivot(index='max_depth',columns='n_estimators',values='accuracy')
ac_df_pca=ac_df_pca.pivot(index='max_depth',columns='n_estimators',values='accuracy')
ac_df_lda=ac_df_lda.pivot(index='max_depth',columns='n_estimators',values='accuracy')

#Plotting the graph
fig, ax=plt.subplots(1,3,figsize=(20, 5), sharey='row')

sns.heatmap(ac_df,annot=True, ax=ax[0])
sns.heatmap(ac_df_pca,annot=True ,ax=ax[1])
sns.heatmap(ac_df_lda,annot=True ,ax=ax[2])

plt.subplots_adjust(wspace=0.40, hspace=0.1)
fig.suptitle('Accuracy Heatplot', fontsize=12)
plt.show()
```



## Gradient Boosting Classifier

### 5-Fold CV and hyperparameter tuning

In [55]:

```

gbc = GradientBoostingClassifier()
gbc.fit(X_train, Y_train)
Y_pred = gbc.predict(X_test)
cm=confusion_matrix(Y_test, Y_pred)
#print(confusion_matrix(Y_test, Y_pred))
#print(classification_report(Y_test, Y_pred))
print("Accuracy Score for Gradient Boosting classifier without hyper
parameter tuning is: ",accuracy_score(Y_test, Y_pred))
#disp =
ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=dtc.classes_)
#disp.plot()

grid_params = {
    'n_estimators': [2,120,200,500,800],
    'learning_rate' : [0.01, 0.1,1],
    #'criterion' : [None,'mse','mae'],
    #'max_depth' : range(2,22,2),
    #'max_features': range(3,10,1),
    #'min_samples_split' : range(2, 10, 1),
    #'min_samples_leaf' : range(2, 10, 1)
}

grid_search = GridSearchCV(gbc, grid_params, cv = 5, n_jobs = -1, verbose =
True)
#print(grid_search.get_params().keys())
grid_search.fit(X_train, Y_train)
gbc = grid_search.best_estimator_
Y_pred = gbc.predict(X_test)
print("Accuracy Score is :",accuracy_score(Y_test, Y_pred))
gbc_acc=accuracy_score(Y_test, Y_pred)
# best parameters and best score
print("Best Parameters are : ",grid_search.best_params_)
gbc_set=grid_search.best_params_
#print("Best Score using PCA is : ",grid_search.best_score_)
gbc_ac_df=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
gbc_ac_df["accuracy"]=grid_search.cv_results_['mean_test_score']

grid_search.fit(X_train_pca, Y_train_pca)
gbc = grid_search.best_estimator_
Y_pred_pca = gbc.predict(X_test_pca)
print("Accuracy Score is :",accuracy_score(Y_test_pca, Y_pred_pca))

```

```

gbc_pca_acc=accuracy_score(Y_test_pca, Y_pred_pca)
# best parameters and best score
print("Best Parameters are : ",grid_search.best_params_)
gbc_pca_set=grid_search.best_params_
#print("Best Score using PCA is : ",grid_search.best_score_)
gbc_ac_df_pca=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
gbc_ac_df_pca["accuracy"]=grid_search.cv_results_['mean_test_score']

grid_search.fit(X_train_lda, Y_train_lda)
gbc = grid_search.best_estimator_
Y_pred_lda = gbc.predict(X_test_lda)
print("Accuracy Score is :",accuracy_score(Y_test_lda, Y_pred_lda))
gbc_lda_acc=accuracy_score(Y_test_lda, Y_pred_lda)
# best parameters and best score
print("Best Parameters are : ",grid_search.best_params_)
gbc_lda_set=grid_search.best_params_
#print("Best Score using PCA is : ",grid_search.best_score_)
gbc_ac_df_lda=pd.DataFrame(grid_search.cv_results_['params'])
#Creating a data frame with hyperparameters and accuracy
gbc_ac_df_lda["accuracy"]=grid_search.cv_results_['mean_test_score']

```

Accuracy Score for Gradient Boosting classifier without hyper parameter tuning is: 0.2535885167464115

Fitting 5 folds for each of 15 candidates, totalling 75 fits

Accuracy Score is : 0.26913875598086123

Best Parameters are : {'learning\_rate': 0.01, 'n\_estimators': 120}

Fitting 5 folds for each of 15 candidates, totalling 75 fits

Accuracy Score is : 0.26435406698564595

Best Parameters are : {'learning\_rate': 0.01, 'n\_estimators': 200}

Fitting 5 folds for each of 15 candidates, totalling 75 fits

Accuracy Score is : 0.28708133971291866

Best Parameters are : {'learning\_rate': 0.01, 'n\_estimators': 120}

In [56]:

```

cm=confusion_matrix(Y_test, Y_pred)
cm_pca=confusion_matrix(Y_test_pca, Y_pred_pca)
cm_lda=confusion_matrix(Y_test_lda, Y_pred_lda)
cm_dict = {
    "raw": cm,
    "pca": cm_pca,
    "lda": cm_lda
}

f, axes = plt.subplots(1, 3, figsize=(20, 5), sharey='row')
for i, (c, cm) in enumerate(cm_dict.items()):
    #print(i, cm)

```

```

disp = ConfusionMatrixDisplay(cm)

disp.plot(ax=axes[i], xticks_rotation=45)

disp.ax_.set_title(c)

disp.im_.colorbar.remove()

disp.ax_.set_xlabel('')

if i!=0:

    disp.ax_.set_ylabel('')


f.text(0.4, 0.1, 'Predicted label', ha='left')

plt.subplots_adjust(wspace=0.40, hspace=0.1)

f.colorbar(disp.im_, ax=axes)

plt.show()

```



## Heatplot of learning\_rate and n\_estimators

In [57]:

```

#Pivoting the dataframe for plotting heat map
gbc_ac_df=gbc_ac_df.pivot(index='learning_rate',columns='n_estimators',values='...

gbc_ac_df_pca=gbc_ac_df_pca.pivot(index='learning_rate',columns='n_estimators',...


gbc_ac_df_lda=gbc_ac_df_lda.pivot(index='learning_rate',columns='n_estimators',...

#Plotting the graph
fig, ax =plt.subplots(1,3,figsize=(20, 5), sharey='row')

sns.heatmap(gbc_ac_df,annot=True, ax=ax[0])
sns.heatmap(gbc_ac_df_pca,annot=True ,ax=ax[1])
sns.heatmap(gbc_ac_df_lda,annot=True ,ax=ax[2])

plt.subplots_adjust(wspace=0.40, hspace=0.1)
fig.suptitle('Accuracy Heatplot', fontsize=12)
plt.show()

```



## Final Results

```
In [66]: dt_dict = {'Algo': ['Decision Tree', 'Random Forest', 'Gradient Boosting'],
                  'abalone_raw': [dt_raw_acc, rf_raw_acc, gbc_acc], 'abalone_pca':
[dt_pca_acc, rf_pca_acc, gbc_pca_acc],
                  'abalone_lda': [dt_lda_acc, rf_lda_acc, gbc_lda_acc]}

df_mid=pd.DataFrame(data=dt_dict, columns=
['Algo', 'abalone_raw', 'abalone_pca', 'abalone_lda'])
#display(df_mid)
df_nb=df_merged.reset_index()
final_df = df_nb.append(df_mid, ignore_index=True)
display(final_df)
```

	Algo	abalone_raw	abalone_pca	abalone_lda
0	KNeighborsClassifier	0.250907	0.249947	0.253777
1	GaussianNB	0.230326	0.232716	0.251631
2	MultinomialNB	0.206377	0.165673	0.164955
3	ComplementNB	0.170705	0.171898	0.228406
4	Decision Tree	0.253589	0.264354	0.234450
5	Random Forest	0.273923	0.267943	0.272727
6	Gradient Boosting	0.269139	0.264354	0.287081

```
In [67]: #print(dt_raw_set.keys())
#print(dt_raw_set.values(),dt_pca_set.values(),dt_lda_set.values())
dt_s =[dt_raw_set,dt_pca_set,dt_lda_set]
df_dt_settings = pd.DataFrame(dt_s)
df_dt_settings['Algo']= 'Decision Tree'
df_dt_settings['Data'] = ['raw','pca','lda']
display(df_dt_settings)
```

	max_depth	min_samples_leaf	min_samples_split	Algo	Data
0	4	2	2	Decision Tree	raw

	max_depth	min_samples_leaf	min_samples_split	Algo	Data
1	4	4	2	Decision Tree	pca
2	3	2	2	Decision Tree	lda

In [68]:

```
rf_s =[rf_raw_set,rf_pca_set,rf_lda_set]
df_rf_settings = pd.DataFrame(rf_s)
df_rf_settings['Algo']= 'Random Forest'
df_rf_settings['Data'] = ['raw','pca','lda']
display(df_rf_settings)
```

	criterion	max_depth	n_estimators	Algo	Data
0	entropy	5	5	Random Forest	raw
1	entropy	6	200	Random Forest	pca
2	entropy	4	500	Random Forest	lda

In [69]:

```
gbc_s =[gbc_set,gbc_pca_set,gbc_lda_set]
df_gbc_settings = pd.DataFrame(gbc_s)
df_gbc_settings['Algo']= 'Gradient Boosting'
df_gbc_settings['Data'] = ['raw','pca','lda']
display(df_gbc_settings)
```

	learning_rate	n_estimators	Algo	Data
0	0.01	120	Gradient Boosting	raw
1	0.01	200	Gradient Boosting	pca
2	0.01	120	Gradient Boosting	lda

In [2]:

```
knn_setting = ({'n_neighbors':66 , 'metric': 'minkowski' , 'weights':
"uniform",'p':2})
gnb_settings = ("default")
mnb_settings = ("default")
cnb_settings = ("default")
dt_settings =
({'criterion':'Entropy','max_depth':4,'min_samples_leaf':4,'min_samples_split':

rf_settings = ({'criterion':'Entropy','max_depth':5,'n_estimators':500})
gb_settings = ({'learning_rate':0.01,'n_estimators':120})

settings_all =
[knn_setting,gnb_settings,mnb_settings,cnb_settings,dt_settings,dt_settings,rf

settings_all
df_final['settings'] =settings_all
```

```
In [5]: display(df_final)
```

```
Out[5]:
```

	Algo	settings	abalone_raw	abalone_pca	abalone_lda
0	KNeighborsClassifier	{'n_neighbors':66, 'metric': 'minkowski', 'w...	0.250907	0.249947	0.253777
1	GaussianNB	default	0.230326	0.232716	0.251631
2	MultinomialNB	default	0.206377	0.165673	0.164955
3	ComplementNB	default	0.170705	0.171898	0.228406
4	Decision Tree	{'criterion':'Entropy','max_depth':4,'min_samp...	0.253589	0.264354	0.234450
5	Random Forest	{'criterion':'Entropy','max_depth':5,'n_estima...	0.273923	0.267943	0.272727
6	Gradient Boosting	{'learning_rate':0.01,'n_estimators':120}	0.269139	0.264354	0.287081

## Comments

- Gradient Boosting performed best with the pipeline (abalone -> scaling -> lda -> scaling -> Gradient Boosting Classifier)
- Dimensionality Reduction helped in Decision tree, Random Forest and Gradient Boosting algorithms, while PCA gave consistent accuracies with mentioned algorithms, LDA performed well only with Gradient Boosting and Random Forest algorithms.
- Decision tree and Random Forest gave better accuracy with lesser max depth as compared to wine dataset.

## References

- [1] P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis. Modeling wine preferences by data mining from physicochemical properties. In Decision Support Systems, Elsevier, 47(4):547-553. ISSN: 0167-9236.
- [2] Dua, D. and Graff, C. (2019). UCI Machine Learning Repository [<http://archive.ics.uci.edu/ml>]. Irvine, CA: University of California, School of Information and Computer Science.
- [3] Compare the effect of different scalers on data with outliers — scikit-learn 0.22.1 documentation. (n.d.). Retrieved January, 2022, from [https://scikitlearn.org/stable/auto\\_examples/preprocessing/plot\\_all\\_scaling.html#sphx-glr-auto-examples-preprocessing-plot-all-scaling-py](https://scikitlearn.org/stable/auto_examples/preprocessing/plot_all_scaling.html#sphx-glr-auto-examples-preprocessing-plot-all-scaling-py) contributed by Scikit-learn: Machine Learning in Python, Pedregosa et al., JMLR 12, pp. 2825-2830, 2011.
- [4] Discussed KNeighbor Classifiers — scikit-learn 0.22.1 documentation. (n.d.). Retrieved January 27, 2022, from <https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html> contributed by Scikit-learn: Machine Learning in Python, Pedregosa et al., JMLR 12, pp. 2825-2830, 2011.
- [5] Discussed Normalization techniques and parameters — scikit-learn 0.22.1 documentation. (n.d.). Retrieved January 27, 2022, from <https://scikit-learn.org/stable/modules/preprocessing.html#normalization> contributed by Scikit-learn: Machine Learning in Python, Pedregosa et al., JMLR 12, pp. 2825-2830, 2011.
- [6] Discussed different distance metrics — scikit-learn 0.22.1 documentation. (n.d.). Retrieved January 27, 2022, from <https://scikit-learn.org/stable/modules/generated/sklearn.metrics.DistanceMetric.html#sklearn.metrics.DistanceMetric>



[7] Discussed axes, plotting and figures in python — matplotlib.org documentation. (n.d.). Retrieved January 27, 2022, from [https://matplotlib.org/stable/api/legend\\_api.html?highlight=legends](https://matplotlib.org/stable/api/legend_api.html?highlight=legends) contributed by J. D. Hunter, "Matplotlib: A 2D Graphics Environment", Computing in Science & Engineering, vol. 9, no. 3, pp. 90-95, 2007.

[8] Discussed plotting of dataframe columns inpython — pandas.org documentation. (n.d.). Retrieved January 27, 2022, from <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.plot.scatter.html> contributed by McKinney, W., & others. (2010). Data structures for statistical computing in python. In Proceedings of the 9th Python in Science Conference (Vol. 445, pp. 51–56).

[9] Singh, S. (2022). Analysis-pca-red-wine-quality-prediction-87%. Retrieved from <https://www.kaggle.com/sonalisingh1411/analysis-pca-red-wine-quality-prediction-87>

[10] Brownlee, J. (2019). Ordinal and One-Hot Encodings for Categorical Data. Retieved from <https://machinelearningmastery.com/>

[10] Katara, V. (2020). Dimensionality Reduction — PCA vs LDA vs t-SNE. Retieved from <https://medium.com/analytics-vidhya/dimensionality-reduction-pca-vs-lda-vs-t-sne-681636bc686>

[11] Kolassa, S. (2020). Setting-max-depth-greater-than-the-number-of-features-in-a-randomforest. Retieved from <https://stats.stackexchange.com/questions/336382/setting-max-depth-greater-than-the-number-of-features-in-a-randomforest#:~:text=There%20is%20no%20problem%20with,two%20features%2C%20Age%20and%20Sex%20.>