

# Assignment 3

## Fashion MNIST Dataset

Date : 01 April 2022

### Loadind Data

In [31]:

```
#importing libraries
import time
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sb

from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.cluster import DBSCAN
from sklearn.neighbors import NearestNeighbors
from sklearn.manifold import TSNE

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.utils import to_categorical

from tensorflow.keras.models import Sequential, load_model, Model
from tensorflow.keras.layers import Dense, Dropout, Flatten, Lambda, Input,
Activation, Conv2D, MaxPooling2D, BatchNormalization, GlobalAveragePooling2D,
Add , TimeDistributed, LSTM
from tensorflow.keras.optimizers import Adam, SGD, Adagrad, Adadelta, RMSprop
#from tensorflow.keras.callbacks import ReduceLROnPlateau,
LearningRateScheduler
#from tensorflow.keras.layers import LeakyReLU
#from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
```

In [32]:

```
#Constants
BATCH_SIZE = 500
EPOCH = 10
n_components = 2
random_state = 27
```

```
In [33]: #loading files into dataframe from mounted google drive
df_x_train = pd.read_csv("/content/drive/MyDrive/ece657a-1221-asg3-
fashionmnist-datafiles/x_train.csv")
df_y_train = pd.read_csv("/content/drive/MyDrive/ece657a-1221-asg3-
fashionmnist-datafiles/y_train.csv")
df_x_test = pd.read_csv("/content/drive/MyDrive/ece657a-1221-asg3-
fashionmnist-datafiles/x_test.csv")
df_y_test = pd.read_csv("/content/drive/MyDrive/ece657a-1221-asg3-
fashionmnist-datafiles/y_test.csv")
```

```
In [34]: #shapes of train and test dataset
print(df_x_test.shape,df_y_test.shape)
print(df_x_train.shape,df_y_train.shape)
```


```
(10000, 784) (10000, 1)
(60000, 784) (60000, 1)
```

```
In [35]: #number of unique classes and their class wise count
classes = np.unique(df_y_train)
nClasses = len(classes)
print('Total number of outputs : ', nClasses)
print('Output classes : ', classes)

plt.bar(df_y_train['0'].unique(),df_y_train.value_counts())
```

```
Total number of outputs : 5
Output classes : [0 1 2 3 4]
<BarContainer object of 5 artists>
```

Out[35]:



## Data Preprocessing

- Dividing training data into training and validation dataset in ratio of 80:20
- Reshaping features in a 28\*28 array and 1 channel(Grayscale)
- Scaling features(x) in (0,1) by dividing by 255 [pixel range : 0-256]
- Converting label into categorical variable [5 classes]

```
In [36]: #splitting df_train in training and validation dataset
x_train, x_val, y_train, y_val = train_test_split(df_x_train, df_y_train,
test_size=0.2, random_state=random_state)
print("Shape of train, validation and test dataset are :
",x_train.shape,x_val.shape, df_x_test.shape)

x_train = x_train.values.reshape((-1, 28, 28, 1))
x_val = x_val.values.reshape((-1, 28, 28, 1))
x_test = df_x_test.values.reshape((-1, 28, 28, 1))

x_train = x_train.astype("float32")/255
x_val = x_val.astype("float32")/255
x_test = x_test.astype("float32")/255
print("Shape of train, validation and test dataset are :
",x_train.shape,x_val.shape,x_test.shape)

y_train = to_categorical(y_train, num_classes=5)
y_val = to_categorical(y_val, num_classes=5)
y_test = to_categorical(df_y_test, num_classes=5)
print("Shape of train, validation and test for dependent variable are :
",y_train.shape,y_val.shape,y_test.shape)

Shape of train, validation and test dataset are : (48000, 784) (12000, 784) (10000, 784)
Shape of train, validation and test dataset are : (48000, 28, 28, 1) (12000, 28, 28, 1)
(10000, 28, 28, 1)
Shape of train, validation and test for dependent variable are : (48000, 5) (12000, 5) (10000, 5)
```

```
In [ ]: #Initial approach for creating validation dataset but scikit in-bulit method
(train_test_split)
#provided a nice implementation and hence skipped this approach
#x_val = x_train[-10000:]
#y_val = y_train[-10000:]
#x_train = x_train[:-10000]
#y_train = y_train[:-10000]
```

## CM1 : Default CNN

### Convolution Neural Network

```
In [37]: #CM1 Model as defined in A3
model= Sequential(name="CM1")
model.add(Conv2D(32, kernel_size=(3, 3),
strides=(1, 1),
activation='relu',
```

```

        input_shape=(28,28,1), padding = 'same', name = 'Conv1'))
model.add(MaxPooling2D(pool_size=(2, 2), name = 'MaxPool'))
model.add(Conv2D(32, (3, 3), activation='relu', name = 'Conv2'))
model.add(Flatten(name='Flatten'))
model.add(Dense(256, activation='relu', name = 'Dense'))
model.add(Dense(5, activation='softmax', name = 'Softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=
['accuracy'])

```

In [38]:

```
model.summary()
```

Model: "CM1"

Layer (type)	Output Shape	Param #
Conv1 (Conv2D)	(None, 28, 28, 32)	320
MaxPool (MaxPooling2D)	(None, 14, 14, 32)	0
Conv2 (Conv2D)	(None, 12, 12, 32)	9248
Flatten (Flatten)	(None, 4608)	0
Dense (Dense)	(None, 256)	1179904
Softmax (Dense)	(None, 5)	1285

=====  
 Total params: 1,190,757  
 Trainable params: 1,190,757  
 Non-trainable params: 0

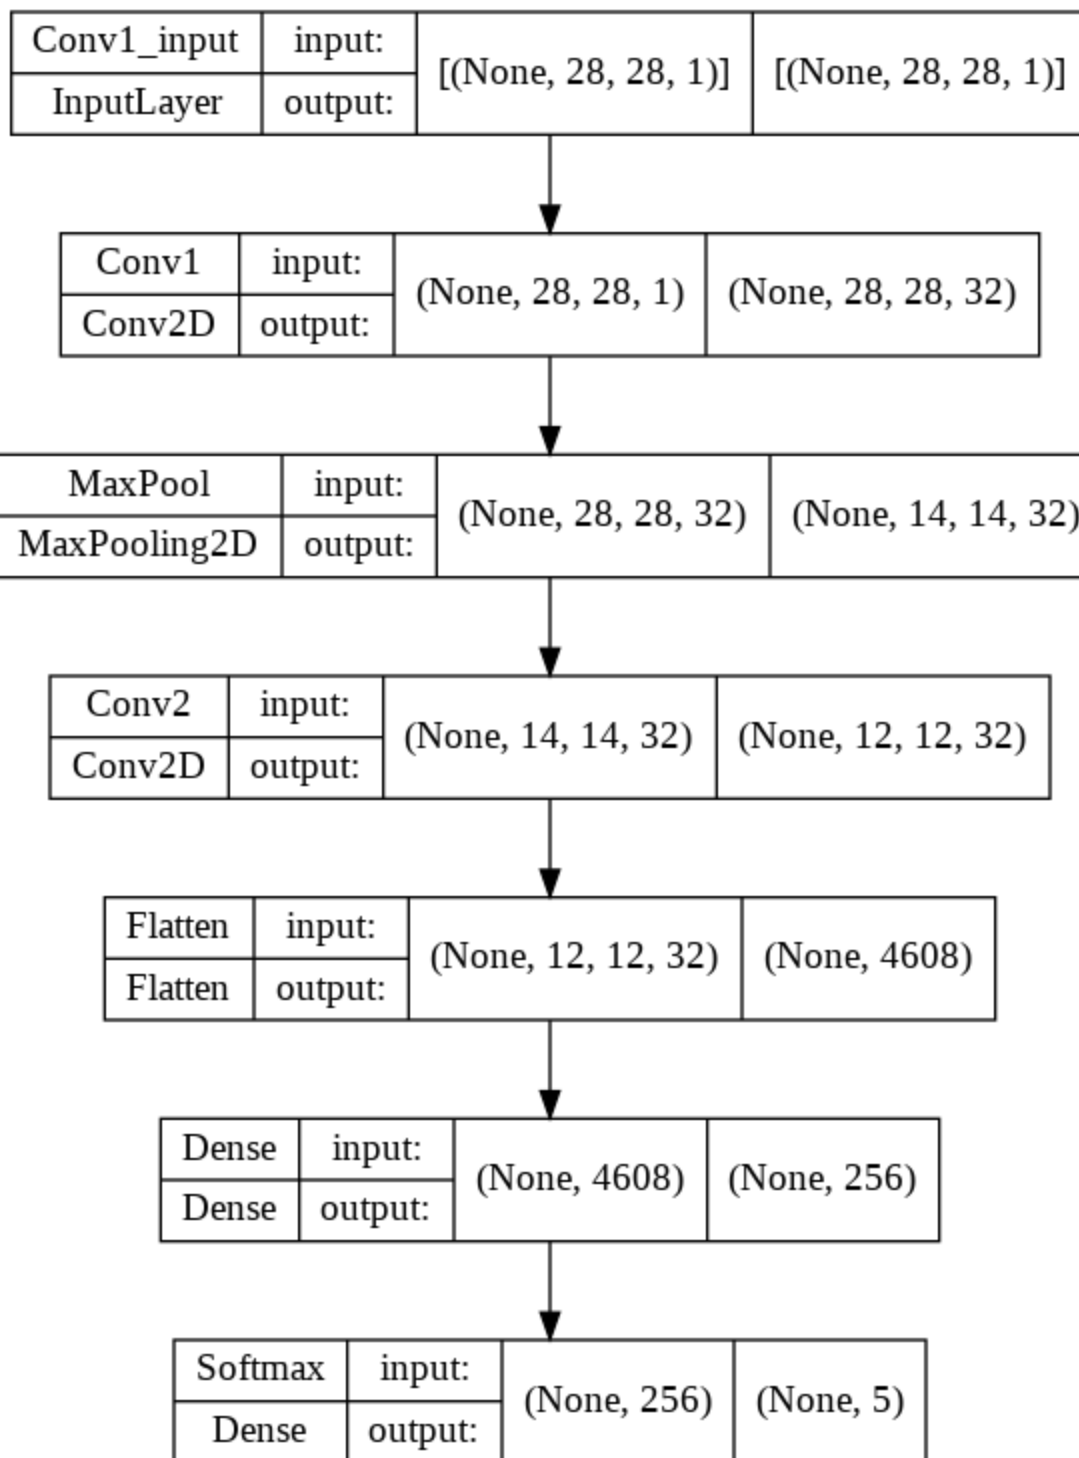
In [39]:

```

#Just another wway to visualize model architecture
tf.keras.utils.plot_model(model, to_file='CM1.png', show_shapes=True)

```

Out[39]:



In [40]:

```
#To include early stopping (depends on patience or loss)
#es = EarlyStopping(monitor='val_loss', mode='min', verbose=1)
#Time performance for train dataset
start = time.time()
#Fitting the model
cm1_history = model.fit(x_train, y_train,
                        batch_size=BATCH_SIZE,
                        epochs=50,
                        shuffle = True,
                        #callbacks=[es],
                        validation_data=(x_val, y_val))
print("Total time: ", time.time() - start, "seconds")
```

Epoch 1/50  
96/96 [=====] - 42s 429ms/step - loss: 0.4832 - accuracy: 0.8102  
- val\_loss: 0.3401 - val\_accuracy: 0.8763  
Epoch 2/50  
96/96 [=====] - 41s 424ms/step - loss: 0.2926 - accuracy: 0.8950  
- val\_loss: 0.2787 - val\_accuracy: 0.9006  
Epoch 3/50  
96/96 [=====] - 41s 422ms/step - loss: 0.2504 - accuracy: 0.9103  
- val\_loss: 0.2400 - val\_accuracy: 0.9154  
Epoch 4/50  
96/96 [=====] - 40s 414ms/step - loss: 0.2158 - accuracy: 0.9241  
- val\_loss: 0.2148 - val\_accuracy: 0.9236  
Epoch 5/50  
96/96 [=====] - 40s 416ms/step - loss: 0.1947 - accuracy: 0.9304  
- val\_loss: 0.2119 - val\_accuracy: 0.9273  
Epoch 6/50  
96/96 [=====] - 39s 409ms/step - loss: 0.1769 - accuracy: 0.9369  
- val\_loss: 0.2041 - val\_accuracy: 0.9258  
Epoch 7/50  
96/96 [=====] - 40s 413ms/step - loss: 0.1656 - accuracy: 0.9408  
- val\_loss: 0.2022 - val\_accuracy: 0.9288  
Epoch 8/50  
96/96 [=====] - 41s 424ms/step - loss: 0.1532 - accuracy: 0.9454  
- val\_loss: 0.2008 - val\_accuracy: 0.9277  
Epoch 9/50  
96/96 [=====] - 40s 420ms/step - loss: 0.1420 - accuracy: 0.9486  
- val\_loss: 0.1834 - val\_accuracy: 0.9345  
Epoch 10/50  
96/96 [=====] - 40s 415ms/step - loss: 0.1349 - accuracy: 0.9516  
- val\_loss: 0.1844 - val\_accuracy: 0.9371  
Epoch 11/50  
96/96 [=====] - 40s 422ms/step - loss: 0.1241 - accuracy: 0.9554  
- val\_loss: 0.1891 - val\_accuracy: 0.9343  
Epoch 12/50  
96/96 [=====] - 41s 425ms/step - loss: 0.1133 - accuracy: 0.9605  
- val\_loss: 0.1799 - val\_accuracy: 0.9350  
Epoch 13/50  
96/96 [=====] - 41s 429ms/step - loss: 0.1093 - accuracy: 0.9607  
- val\_loss: 0.1776 - val\_accuracy: 0.9392  
Epoch 14/50  
96/96 [=====] - 41s 423ms/step - loss: 0.0968 - accuracy: 0.9661  
- val\_loss: 0.1746 - val\_accuracy: 0.9389  
Epoch 15/50  
96/96 [=====] - 40s 421ms/step - loss: 0.0895 - accuracy: 0.9696  
- val\_loss: 0.1782 - val\_accuracy: 0.9407  
Epoch 16/50  
96/96 [=====] - 40s 422ms/step - loss: 0.0849 - accuracy: 0.9708  
- val\_loss: 0.1917 - val\_accuracy: 0.9382  
Epoch 17/50  
96/96 [=====] - 40s 418ms/step - loss: 0.0785 - accuracy: 0.9727  
- val\_loss: 0.1868 - val\_accuracy: 0.9389  
Epoch 18/50  
96/96 [=====] - 40s 419ms/step - loss: 0.0722 - accuracy: 0.9748  
- val\_loss: 0.1917 - val\_accuracy: 0.9391  
Epoch 19/50  
96/96 [=====] - 40s 413ms/step - loss: 0.0686 - accuracy: 0.9763  
- val\_loss: 0.1970 - val\_accuracy: 0.9377  
Epoch 20/50

```
96/96 [=====] - 40s 415ms/step - loss: 0.0612 - accuracy: 0.9792
- val_loss: 0.2050 - val_accuracy: 0.9373
Epoch 21/50
96/96 [=====] - 40s 417ms/step - loss: 0.0560 - accuracy: 0.9810
- val_loss: 0.2050 - val_accuracy: 0.9386
Epoch 22/50
96/96 [=====] - 40s 417ms/step - loss: 0.0507 - accuracy: 0.9831
- val_loss: 0.2064 - val_accuracy: 0.9391
Epoch 23/50
96/96 [=====] - 40s 415ms/step - loss: 0.0485 - accuracy: 0.9840
- val_loss: 0.2147 - val_accuracy: 0.9380
Epoch 24/50
96/96 [=====] - 41s 424ms/step - loss: 0.0434 - accuracy: 0.9857
- val_loss: 0.1990 - val_accuracy: 0.9409
Epoch 25/50
96/96 [=====] - 40s 418ms/step - loss: 0.0396 - accuracy: 0.9870
- val_loss: 0.2063 - val_accuracy: 0.9404
Epoch 26/50
96/96 [=====] - 40s 413ms/step - loss: 0.0318 - accuracy: 0.9904
- val_loss: 0.2218 - val_accuracy: 0.9408
Epoch 27/50
96/96 [=====] - 40s 416ms/step - loss: 0.0306 - accuracy: 0.9900
- val_loss: 0.2193 - val_accuracy: 0.9434
Epoch 28/50
96/96 [=====] - 40s 417ms/step - loss: 0.0264 - accuracy: 0.9920
- val_loss: 0.2381 - val_accuracy: 0.9436
Epoch 29/50
96/96 [=====] - 40s 419ms/step - loss: 0.0242 - accuracy: 0.9927
- val_loss: 0.2285 - val_accuracy: 0.9418
Epoch 30/50
96/96 [=====] - 40s 417ms/step - loss: 0.0217 - accuracy: 0.9934
- val_loss: 0.2387 - val_accuracy: 0.9413
Epoch 31/50
96/96 [=====] - 47s 486ms/step - loss: 0.0174 - accuracy: 0.9950
- val_loss: 0.2518 - val_accuracy: 0.9414
Epoch 32/50
96/96 [=====] - 40s 415ms/step - loss: 0.0145 - accuracy: 0.9961
- val_loss: 0.2528 - val_accuracy: 0.9433
Epoch 33/50
96/96 [=====] - 40s 414ms/step - loss: 0.0132 - accuracy: 0.9967
- val_loss: 0.2634 - val_accuracy: 0.9394
Epoch 34/50
96/96 [=====] - 40s 415ms/step - loss: 0.0142 - accuracy: 0.9959
- val_loss: 0.2727 - val_accuracy: 0.9404
Epoch 35/50
96/96 [=====] - 40s 414ms/step - loss: 0.0111 - accuracy: 0.9973
- val_loss: 0.3108 - val_accuracy: 0.9401
Epoch 36/50
96/96 [=====] - 40s 415ms/step - loss: 0.0131 - accuracy: 0.9962
- val_loss: 0.2890 - val_accuracy: 0.9416
Epoch 37/50
96/96 [=====] - 40s 416ms/step - loss: 0.0068 - accuracy: 0.9986
- val_loss: 0.2889 - val_accuracy: 0.9413
Epoch 38/50
96/96 [=====] - 40s 413ms/step - loss: 0.0060 - accuracy: 0.9988
- val_loss: 0.2919 - val_accuracy: 0.9417
Epoch 39/50
96/96 [=====] - 40s 413ms/step - loss: 0.0036 - accuracy: 0.9997
```

```

- val_loss: 0.3074 - val_accuracy: 0.9421
Epoch 40/50
96/96 [=====] - 40s 413ms/step - loss: 0.0034 - accuracy: 0.9996
- val_loss: 0.3106 - val_accuracy: 0.9414
Epoch 41/50
96/96 [=====] - 40s 419ms/step - loss: 0.0039 - accuracy: 0.9994
- val_loss: 0.3080 - val_accuracy: 0.9441
Epoch 42/50
96/96 [=====] - 40s 419ms/step - loss: 0.0030 - accuracy: 0.9996
- val_loss: 0.3146 - val_accuracy: 0.9427
Epoch 43/50
96/96 [=====] - 40s 414ms/step - loss: 0.0048 - accuracy: 0.9988
- val_loss: 0.3757 - val_accuracy: 0.9327
Epoch 44/50
96/96 [=====] - 40s 412ms/step - loss: 0.0346 - accuracy: 0.9873
- val_loss: 0.2917 - val_accuracy: 0.9358
Epoch 45/50
96/96 [=====] - 40s 412ms/step - loss: 0.0145 - accuracy: 0.9950
- val_loss: 0.3042 - val_accuracy: 0.9414
Epoch 46/50
96/96 [=====] - 40s 414ms/step - loss: 0.0060 - accuracy: 0.9986
- val_loss: 0.3177 - val_accuracy: 0.9436
Epoch 47/50
96/96 [=====] - 40s 417ms/step - loss: 0.0019 - accuracy: 0.9999
- val_loss: 0.3296 - val_accuracy: 0.9447
Epoch 48/50
96/96 [=====] - 40s 418ms/step - loss: 9.5341e-04 - accuracy: 1.0
000 - val_loss: 0.3394 - val_accuracy: 0.9437
Epoch 49/50
96/96 [=====] - 40s 418ms/step - loss: 8.4584e-04 - accuracy: 1.0
000 - val_loss: 0.3445 - val_accuracy: 0.9446
Epoch 50/50
96/96 [=====] - 40s 414ms/step - loss: 6.3437e-04 - accuracy: 1.0
000 - val_loss: 0.3502 - val_accuracy: 0.9449
Total time: 2063.238160133362 seconds

```

In [41]:

```

#evaluation
_, train_acc = model.evaluate(x_train, y_train, verbose=0)
_, val_acc = model.evaluate(x_val, y_val, verbose=0)
_, test_acc = model.evaluate(x_test, y_test, verbose=0)
print('Train: %.3f, Val: %.3f,Test: %.3f' % (train_acc*100, val_acc*100,
test_acc*100))

```

Train: 100.000, Val: 94.492,Test: 93.960

In [42]:

```

#Time performance for test dataset
start_time = time.time()
y_pred = model.predict(x_test)
end_time = time.time()
time_taken = end_time - start_time
print("Testing time : ", time_taken)

```

Testing time : 3.376964807510376

In [43]:



*#Plots: Visualization of accuracy vs epoch and loss vs epoch*


```
plt.plot(cm1_history.history['accuracy'])
plt.plot(cm1_history.history['val_accuracy'])
plt.title('Model accuracy vs Epoch')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='lower right')
plt.show()

# summarize history for loss
plt.plot(cm1_history.history['loss'])
plt.plot(cm1_history.history['val_loss'])
plt.title('Model loss vs Epoch')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper right')
plt.show()


plt.plot(cm1_history.history['val_accuracy'])
plt.plot(cm1_history.history['val_loss'])
plt.title('Validation Accuracy vs Validation Loss based on Epochs')
plt.ylabel('Validation Accuracy, Test Loss')
plt.xlabel('Epochs')
plt.legend(['accuracy', 'loss'], loc='upper right')
plt.show()

plt.plot(cm1_history.history['accuracy'])
plt.plot(cm1_history.history['loss'])
plt.title('Training Accuracy vs Training Loss based on Epochs')
plt.ylabel('Train Accuracy, Train Loss')
plt.xlabel('Epochs')
plt.legend(['accuracy', 'loss'], loc='upper right')
plt.show()
```


Model accuracy vs Epoch




Model loss vs Epoch



Validation Accuracy vs Validation Loss based on Epochs





### Comments:

- We obtain ~94% accuracy on test dataset for this simple CNN model.
- We tried with different epochs to see the impact of local minima and observed that validation loss started increasing and validation accuracy fluctuated in small increments but on both sides of scale. The data gave significant accuracy in few epochs (as low as 5).
- Adam Optimizer yielded better accuracy compared to SGD.
- Categorical\_crossentropy loss function gave better accuracy measure compared to other loss functions.
- Runtime for training and testing suggests good performance (quite fast).

## CM2 : Own Network

We implement a simple CNN model, a variant of CM1 model by adding a conv layer and dense layer. We implement more sophisticated CNN architectures in subsequent sections to compare the impact of different components in architecture.

### CNN1

**To experiment with the impact of iterative increase and decrease in filter size (increasing and decreasing volume) and an added dense layer (better relationship in features)**

- We use Adaptive Moment Estimation (Adam) optimizer with default value
- We make use of categorical cross entropy as loss function
- We make use of 3 Conv2D layer but with different kernel sizes (3, 3 and 3 respectively) followed by a flatten layer and 2 Dense layers with 256 and 128 neurons and finally a softmax max function to give probabilities in (0,1) range for our 5 class classification task. Complete model architecture is given below.
- We make use of ReLU activation function (Leaky ReLU yielded similar performance)
- Model converged faster with Adam Optimizer than SGD.

```
In [ ]: model = Sequential(name="CNN1")
model.add(Conv2D(32, kernel_size=(3, 3), strides=(1, 1), activation='relu',
                input_shape=(28, 28, 1), padding = 'same', name = 'Conv1'))
model.add(Conv2D(64, (3, 3), activation='relu', name = 'Conv2'))
model.add(MaxPooling2D(pool_size=(2, 2), name = 'MaxPool'))
model.add(Conv2D(32, (3, 3), activation='relu', name = 'Conv3'))
```

```

model.add(Flatten(name='Flatten'))
model.add(Dense(256, activation='relu', name = 'Dense1'))
model.add(Dense(128, activation='relu', name = 'Dense2'))
model.add(Dense(5, activation='softmax', name = 'Softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=
['accuracy'])

```

```
In [ ]: model.summary()
```

Model: "CNN1"

Layer (type)	Output Shape	Param #
=====		
Conv1 (Conv2D)	(None, 28, 28, 32)	320
Conv2 (Conv2D)	(None, 26, 26, 64)	18496
MaxPool (MaxPooling2D)	(None, 13, 13, 64)	0
Conv3 (Conv2D)	(None, 11, 11, 32)	18464
Flatten (Flatten)	(None, 3872)	0
Dense1 (Dense)	(None, 256)	991488
Dense2 (Dense)	(None, 128)	32896
Softmax (Dense)	(None, 5)	645
=====		
Total params: 1,062,309		
Trainable params: 1,062,309		
Non-trainable params: 0		

```

In [ ]: start = time.time()
CNN1_history = model.fit(x_train, y_train,
                        batch_size=BATCH_SIZE,
                        epochs=EPOCH,
                        shuffle = True,
                        validation_data=(x_val, y_val))
print("Total time: ", time.time() - start, "seconds")

```

Epoch 1/10

96/96 [=====] - 152s 2s/step - loss: 0.5250 - accuracy: 0.7914 - val\_loss: 0.3352 - val\_accuracy: 0.8779

Epoch 2/10

96/96 [=====] - 120s 1s/step - loss: 0.2822 - accuracy: 0.8986 - val\_loss: 0.2608 - val\_accuracy: 0.9069

Epoch 3/10

96/96 [=====] - 120s 1s/step - loss: 0.2262 - accuracy: 0.9187 - val\_loss: 0.2260 - val\_accuracy: 0.9189

Epoch 4/10

```
96/96 [=====] - 119s 1s/step - loss: 0.1879 - accuracy: 0.9334 -  
val_loss: 0.2146 - val_accuracy: 0.9225  
Epoch 5/10  
96/96 [=====] - 119s 1s/step - loss: 0.1669 - accuracy: 0.9395 -  
val_loss: 0.1893 - val_accuracy: 0.9324  
Epoch 6/10  
96/96 [=====] - 118s 1s/step - loss: 0.1508 - accuracy: 0.9456 -  
val_loss: 0.1865 - val_accuracy: 0.9327  
Epoch 7/10  
96/96 [=====] - 116s 1s/step - loss: 0.1307 - accuracy: 0.9534 -  
val_loss: 0.1840 - val_accuracy: 0.9328  
Epoch 8/10  
96/96 [=====] - 117s 1s/step - loss: 0.1179 - accuracy: 0.9566 -  
val_loss: 0.1804 - val_accuracy: 0.9387  
Epoch 9/10  
96/96 [=====] - 115s 1s/step - loss: 0.0996 - accuracy: 0.9638 -  
val_loss: 0.1665 - val_accuracy: 0.9426  
Epoch 10/10  
96/96 [=====] - 115s 1s/step - loss: 0.0896 - accuracy: 0.9679 -  
val_loss: 0.1828 - val_accuracy: 0.9397  
Total time: 1222.489366531372 seconds
```

```
In [ ]: #evaluation  
_, train_acc = model.evaluate(x_train, y_train, verbose=0)  
_, val_acc = model.evaluate(x_val, y_val, verbose=0)  
_, test_acc = model.evaluate(x_test, y_test, verbose=0)  
print('Train: %.3f, Val: %.3f,Test: %.3f' % (train_acc*100, val_acc*100,  
test_acc*100))
```

Train: 97.206, Val: 93.967,Test: 93.820

```
In [ ]: #Time performance for test dataset  
start_time = time.time()  
y_pred = model.predict(x_test)  
end_time = time.time()  
time_taken = end_time - start_time  
print("Testing time : ", time_taken)
```

Testing time : 5.968332529067993

```
In [ ]: #Plots: Visualization of accuracy vs epoch and loss vs epoch  
plt.plot(CNN1_history.history['accuracy'])  
plt.plot(CNN1_history.history['val_accuracy'])  
plt.title('Model accuracy vs Epoch')  
plt.ylabel('accuracy')  
plt.xlabel('epoch')  
plt.legend(['train', 'validation'], loc='lower right')  
plt.show()  
# summarize history for loss  
plt.plot(CNN1_history.history['loss'])  
plt.plot(CNN1_history.history['val_loss'])
```


```




plt.title('Model loss vs Epoch')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper right')
plt.show()

plt.plot(CNN1_history.history['val_accuracy'])
plt.plot(CNN1_history.history['val_loss'])
plt.title('Validation Accuracy vs Validation Loss based on Epochs')
plt.ylabel('Validation Accuracy, Test Loss')
plt.xlabel('Epochs')
plt.legend(['accuracy', 'loss'], loc='upper right')
plt.show()

plt.plot(CNN1_history.history['accuracy'])
plt.plot(CNN1_history.history['loss'])
plt.title('Training Accuracy vs Training Loss based on Epochs')
plt.ylabel('Train Accuracy, Train Loss')
plt.xlabel('Epochs')
plt.legend(['accuracy', 'loss'], loc='upper right')
plt.show()

```





## Comments

We observed ~94% accuracy with our own CNN model.

## CM3 : Own Network

- CNN2 : Adding Dropout layers
- CNN3 : Experimenting impact of Optimizers (SGD)
- CNN4 : Experimentation with activation function (Leaky ReLU)

- CNN\_VGG : Inspiration from VGG16 (focusing on increasing volume)
- LSTM : Long Short term memory architecture

## CNN2

We added Dropout layers to see the impact of regularization. We use rate=0.3 as argument i.e, the fraction of the input units to be dropped at each layer. It is a fraction between 0 and 1 and is used to prevent overfitting

```
In [ ]: model = Sequential(name="CNN_2")

model.add(Conv2D(32, kernel_size=(3, 3),
                 strides=(1, 1),
                 activation='relu',
                 input_shape=(28,28,1), padding = 'same', name = 'Conv1'))
model.add(MaxPooling2D(pool_size=(2, 2), name = 'MaxPool1'))
model.add(Dropout(0.3))

model.add(Conv2D(32, (3, 3), activation='relu', name = 'Conv2'))
model.add(MaxPooling2D(pool_size=(2, 2), name = 'MaxPool2'))
model.add(Dropout(0.3))

model.add(Flatten(name = 'Flatten'))

model.add(Dense(256, activation='relu', name = 'Dense1'))
model.add(Dropout(0.5))

model.add(Dense(64, activation='relu', name = 'Dense2'))

model.add(Dense(5, activation='softmax', name = 'Softmax'))

model.compile(loss='categorical_crossentropy', optimizer='adam',
              metrics=['accuracy'])
```

```
In [ ]: model.summary()
```

Model: "CNN\_2"

Layer (type)	Output Shape	Param #
=====		
Conv1 (Conv2D)	(None, 28, 28, 32)	320
MaxPool1 (MaxPooling2D)	(None, 14, 14, 32)	0
dropout (Dropout)	(None, 14, 14, 32)	0



Conv2 (Conv2D)	(None, 12, 12, 32)	9248
MaxPool2 (MaxPooling2D)	(None, 6, 6, 32)	0
dropout_1 (Dropout)	(None, 6, 6, 32)	0
Flatten (Flatten)	(None, 1152)	0
Dense1 (Dense)	(None, 256)	295168
dropout_2 (Dropout)	(None, 256)	0
Dense2 (Dense)	(None, 64)	16448
Softmax (Dense)	(None, 5)	325

```

=====
Total params: 321,509
Trainable params: 321,509
Non-trainable params: 0

```

In [ ]:

```

start = time.time()
CNN_2_history = model.fit(x_train, y_train,
                          batch_size=BATCH_SIZE,
                          epochs=EPOCH,
                          shuffle = True,
                          validation_data=(x_val, y_val))
print("Total time: ", time.time() - start, "seconds")

```

```

Epoch 1/10
96/96 [=====] - 31s 318ms/step - loss: 0.7341 - accuracy: 0.6894
- val_loss: 0.4416 - val_accuracy: 0.8261
Epoch 2/10
96/96 [=====] - 31s 318ms/step - loss: 0.4404 - accuracy: 0.8289
- val_loss: 0.3444 - val_accuracy: 0.8695
Epoch 3/10
96/96 [=====] - 31s 320ms/step - loss: 0.3754 - accuracy: 0.8585
- val_loss: 0.2996 - val_accuracy: 0.8892
Epoch 4/10
96/96 [=====] - 31s 320ms/step - loss: 0.3337 - accuracy: 0.8760
- val_loss: 0.2678 - val_accuracy: 0.9025
Epoch 5/10
96/96 [=====] - 30s 317ms/step - loss: 0.3094 - accuracy: 0.8867
- val_loss: 0.2470 - val_accuracy: 0.9093
Epoch 6/10
96/96 [=====] - 31s 319ms/step - loss: 0.2866 - accuracy: 0.8951
- val_loss: 0.2342 - val_accuracy: 0.9116
Epoch 7/10
96/96 [=====] - 31s 319ms/step - loss: 0.2725 - accuracy: 0.8998
- val_loss: 0.2223 - val_accuracy: 0.9168
Epoch 8/10
96/96 [=====] - 30s 317ms/step - loss: 0.2577 - accuracy: 0.9076
- val_loss: 0.2118 - val_accuracy: 0.9209
Epoch 9/10
96/96 [=====] - 30s 318ms/step - loss: 0.2476 - accuracy: 0.9086

```

```
- val_loss: 0.2055 - val_accuracy: 0.9248
Epoch 10/10
96/96 [=====] - 30s 317ms/step - loss: 0.2410 - accuracy: 0.9117
- val_loss: 0.2006 - val_accuracy: 0.9268
Total time: 322.5631422996521 seconds
```

```
In [ ]: #evaluation
_, train_acc = model.evaluate(x_train, y_train, verbose=0)
_, val_acc = model.evaluate(x_val, y_val, verbose=0)
_, test_acc = model.evaluate(x_test, y_test, verbose=0)
print('Train: %.3f, Val: %.3f,Test: %.3f' % (train_acc*100, val_acc*100,
test_acc*100))

#Time performance for test dataset
start_time = time.time()
y_pred = model.predict(x_test)
end_time = time.time()
time_taken = end_time - start_time
print("Testing time : ", time_taken)
```




```
Train: 93.448, Val: 92.683,Test: 92.420
Testing time : 1.8049075603485107
```


```
In [ ]: plt.plot(CNN_2_history.history['accuracy'])
plt.plot(CNN_2_history.history['val_accuracy'])
plt.title('Model accuracy vs Epoch')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='lower right')
plt.show()

# summarize history for loss
plt.plot(CNN_2_history.history['loss'])
plt.plot(CNN_2_history.history['val_loss'])
plt.title('Model loss vs Epoch')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper right')
plt.show()

plt.plot(CNN_2_history.history['val_accuracy'])
plt.plot(CNN_2_history.history['val_loss'])
plt.title('Validation Accuracy vs Validation Loss based on Epochs')
plt.ylabel('Validation Accuracy, Test Loss')
plt.xlabel('Epochs')
plt.legend(['accuracy', 'loss'], loc='upper right')
plt.show()
```

```
plt.plot(CNN_2_history.history['accuracy'])
plt.plot(CNN_2_history.history['loss'])
plt.title('Training Accuracy vs Training Loss based on Epochs')
plt.ylabel('Train Accuracy, Train Loss')
plt.xlabel('Epochs')
plt.legend(['accuracy', 'loss'], loc='upper right')
plt.show()
```





## CNN3

Trained the model with Stochastic Gradient Descent (SGD) Optimizer.

```
In [ ]: model = Sequential(name="CNN_3")

model.add(Conv2D(32, kernel_size=(3, 3),
                strides=(1, 1),
                activation='relu',
                input_shape=(28,28,1), padding = 'same', name = 'Conv1'))
model.add(MaxPooling2D(pool_size=(2, 2), name = 'MaxPool1'))
model.add(Conv2D(32, (3, 3), activation='relu', name = 'Conv2'))
model.add(MaxPooling2D(pool_size=(2, 2), name = 'MaxPool2'))
model.add(Dropout(0.3))
model.add(Flatten(name = 'Flatten'))
model.add(Dense(256, activation='relu', name = 'Dense1'))
model.add(Dropout(0.5))
model.add(Dense(64, activation='relu', name = 'Dense2'))
model.add(Dense(5, activation='softmax', name = 'Softmax'))
model.compile(loss='categorical_crossentropy', optimizer='SGD',
              metrics=['accuracy'])
```

```
In [ ]: start = time.time()
CNN_3_history = model.fit(x_train, y_train,
                        batch_size=BATCH_SIZE,
                        epochs=EPOCH,
                        shuffle = True,
                        validation_data=(x_val, y_val))
print("Total time: ", time.time() - start, "seconds")
```

Epoch 1/10

96/96 [=====] - 28s 292ms/step - loss: 1.5011 - accuracy: 0.3079

```

- val_loss: 1.3404 - val_accuracy: 0.5020
Epoch 2/10
96/96 [=====] - 28s 292ms/step - loss: 1.2983 - accuracy: 0.4315
- val_loss: 1.0609 - val_accuracy: 0.6069
Epoch 3/10
96/96 [=====] - 29s 306ms/step - loss: 1.1244 - accuracy: 0.5020
- val_loss: 0.9144 - val_accuracy: 0.6252
Epoch 4/10
96/96 [=====] - 30s 311ms/step - loss: 1.0229 - accuracy: 0.5493
- val_loss: 0.8435 - val_accuracy: 0.6422
Epoch 5/10
96/96 [=====] - 28s 294ms/step - loss: 0.9528 - accuracy: 0.5888
- val_loss: 0.7829 - val_accuracy: 0.6808
Epoch 6/10
96/96 [=====] - 28s 296ms/step - loss: 0.9059 - accuracy: 0.6157
- val_loss: 0.7398 - val_accuracy: 0.7044
Epoch 7/10
96/96 [=====] - 29s 298ms/step - loss: 0.8619 - accuracy: 0.6402
- val_loss: 0.7052 - val_accuracy: 0.7117
Epoch 8/10
96/96 [=====] - 28s 296ms/step - loss: 0.8276 - accuracy: 0.6572
- val_loss: 0.6736 - val_accuracy: 0.7303
Epoch 9/10
96/96 [=====] - 29s 298ms/step - loss: 0.7942 - accuracy: 0.6740
- val_loss: 0.6409 - val_accuracy: 0.7535
Epoch 10/10
96/96 [=====] - 28s 295ms/step - loss: 0.7686 - accuracy: 0.6839
- val_loss: 0.6189 - val_accuracy: 0.7540
Total time: 322.4274568557739 seconds

```

In [ ]:

```

#evaluation
_, train_acc = model.evaluate(x_train, y_train, verbose=0)
_, val_acc = model.evaluate(x_val, y_val, verbose=0)
_, test_acc = model.evaluate(x_test, y_test, verbose=0)
print('Train: %.3f, Val: %.3f,Test: %.3f' % (train_acc*100, val_acc*100,
test_acc*100))

#Time performance for test dataset
start_time = time.time()
y_pred = model.predict(x_test)
end_time = time.time()
time_taken = end_time - start_time
print("Testing time : ", time_taken)

```

```

Train: 75.352, Val: 75.400,Test: 75.210
Testing time : 1.7907054424285889

```

## CNN4

We trained the model with LeakyRelu activation function. It is a variant of ReLU and introduces a small gradient for negative value instead of 0 as in ReLU. Both are similar in that their derivative is monotonic and continuous, They both are able to solve the problem of exploding and vanishing gradients.

```
In [ ]: model = Sequential(name="CNN_4")

model.add(Conv2D(32, kernel_size=(3, 3),
                strides=(1, 1),
                activation=tf.keras.layers.LeakyReLU(alpha=0.2),
                input_shape=(28,28,1), padding = 'same', name = 'Conv1'))
model.add(MaxPooling2D(pool_size=(2, 2), name = 'MaxPool1'))
model.add(Conv2D(32, (3, 3),
                activation=tf.keras.layers.LeakyReLU(alpha=0.2), name = 'Conv2'))
model.add(MaxPooling2D(pool_size=(2, 2), name = 'MaxPool2'))
model.add(Dropout(0.3))
model.add(Flatten(name = 'Flatten'))
model.add(Dense(256, activation=tf.keras.layers.LeakyReLU(alpha=0.2), name =
'Dense1'))
model.add(Dropout(0.5))
model.add(Dense(64, activation=tf.keras.layers.LeakyReLU(alpha=0.2), name =
'Dense2'))
model.add(Dense(5, activation='softmax', name = 'Softmax'))
model.compile(loss='categorical_crossentropy', optimizer='Adam',
              metrics=['accuracy'])
```

```
In [ ]: start = time.time()

CNN_4_history = model.fit(x_train, y_train,
                        batch_size=BATCH_SIZE,
                        epochs=EPOCH,
                        shuffle = True,
                        validation_data=(x_val, y_val))
print("Total time: ", time.time() - start, "seconds")
```

```
Epoch 1/10
96/96 [=====] - 32s 327ms/step - loss: 0.6416 - accuracy: 0.7396
- val_loss: 0.3685 - val_accuracy: 0.8621
Epoch 2/10
96/96 [=====] - 31s 324ms/step - loss: 0.3715 - accuracy: 0.8617
- val_loss: 0.3028 - val_accuracy: 0.8909
Epoch 3/10
96/96 [=====] - 31s 322ms/step - loss: 0.3203 - accuracy: 0.8828
- val_loss: 0.2707 - val_accuracy: 0.9013
Epoch 4/10
96/96 [=====] - 31s 322ms/step - loss: 0.2886 - accuracy: 0.8949
- val_loss: 0.2443 - val_accuracy: 0.9107
Epoch 5/10
96/96 [=====] - 31s 322ms/step - loss: 0.2645 - accuracy: 0.9022
- val_loss: 0.2264 - val_accuracy: 0.9176
Epoch 6/10
96/96 [=====] - 31s 322ms/step - loss: 0.2501 - accuracy: 0.9099
- val_loss: 0.2211 - val_accuracy: 0.9201
Epoch 7/10
```

```

96/96 [=====] - 31s 322ms/step - loss: 0.2366 - accuracy: 0.9143
- val_loss: 0.2096 - val_accuracy: 0.9241
Epoch 8/10
96/96 [=====] - 31s 324ms/step - loss: 0.2230 - accuracy: 0.9204
- val_loss: 0.2012 - val_accuracy: 0.9262
Epoch 9/10
96/96 [=====] - 31s 323ms/step - loss: 0.2147 - accuracy: 0.9211
- val_loss: 0.1912 - val_accuracy: 0.9307
Epoch 10/10
96/96 [=====] - 31s 325ms/step - loss: 0.2042 - accuracy: 0.9262
- val_loss: 0.1843 - val_accuracy: 0.9326
Total time: 322.5065915584564 seconds

```

In [ ]:

```

#evaluation
_, train_acc = model.evaluate(x_train, y_train, verbose=0)
_, val_acc = model.evaluate(x_val, y_val, verbose=0)
_, test_acc = model.evaluate(x_test, y_test, verbose=0)
print('Train: %.3f, Val: %.3f,Test: %.3f' % (train_acc*100, val_acc*100,
test_acc*100))

#Time performance for test dataset
start_time = time.time()
y_pred = model.predict(x_test)
end_time = time.time()
time_taken = end_time - start_time
print("Testing time : ", time_taken)

```

```

Train: 94.356, Val: 93.258,Test: 92.960
Testing time : 2.2175590991973877

```

## CNN\_VGG

VGG is an popular image classification model using fewer CNN layers compared to other image classification algorithms such as AlexNet, LeNet and others. It successively increases the volume by adding CNN layers with a larger filter size and finally using multiple fully connected layers to learn relationship between features.

In [ ]:

```

model = Sequential(name="CNN_VGG")

model.add(Conv2D(32, kernel_size=(3, 3),
                strides=(1, 1),
                activation='relu',
                input_shape=(28,28,1), padding = 'same', name = 'Conv1'))
model.add(MaxPooling2D(pool_size=(2, 2),name = 'MaxPool1'))
model.add(Conv2D(64, (3, 3), activation='relu',name = 'Conv2'))

model.add(Conv2D(128, (3, 3), activation='relu',name = 'Conv3'))
model.add(Conv2D(128, (3, 3), activation='relu',name = 'Conv4'))

```

```

model.add(Conv2D(256, (3, 3), activation='relu', name = 'Conv5'))
model.add(Conv2D(256, (3, 3), activation='relu', name = 'Conv6'))

model.add(MaxPooling2D(pool_size=(2, 2), name = 'MaxPool4'))
model.add(Flatten(name = 'Flatten'))

model.add(Dense(1024, activation='relu', name = 'Dense1'))
model.add(Dropout(0.5))
model.add(Dense(512, activation='relu', name = 'Dense2'))
model.add(Dropout(0.5))
model.add(Dense(256, activation='relu', name = 'Dense3'))
model.add(Dropout(0.5))
model.add(Dense(5, activation='softmax', name = 'Softmax'))

model.compile(loss='categorical_crossentropy', optimizer=
tf.keras.optimizers.Adam(0.003), metrics=['accuracy'])

```

In [ ]: `model.summary()`

Model: "CNN\_VGG"

Layer (type)	Output Shape	Param #
=====		
Conv1 (Conv2D)	(None, 28, 28, 32)	320
MaxPool1 (MaxPooling2D)	(None, 14, 14, 32)	0
Conv2 (Conv2D)	(None, 12, 12, 64)	18496
Conv3 (Conv2D)	(None, 10, 10, 128)	73856
Conv4 (Conv2D)	(None, 8, 8, 128)	147584
Conv5 (Conv2D)	(None, 6, 6, 256)	295168
Conv6 (Conv2D)	(None, 4, 4, 256)	590080
MaxPool4 (MaxPooling2D)	(None, 2, 2, 256)	0
Flatten (Flatten)	(None, 1024)	0
Dense1 (Dense)	(None, 1024)	1049600
dropout_7 (Dropout)	(None, 1024)	0
Dense2 (Dense)	(None, 512)	524800
dropout_8 (Dropout)	(None, 512)	0
Dense3 (Dense)	(None, 256)	131328



dropout_9 (Dropout)	(None, 256)	0
Softmax (Dense)	(None, 5)	1285

```

=====
Total params: 2,832,517
Trainable params: 2,832,517
Non-trainable params: 0

```

```

In [ ]: start = time.time()
CNN_VGG_history = model.fit(x_train, y_train,
                             batch_size=400,
                             epochs=5,
                             shuffle = True,
                             validation_data=(x_val, y_val))
print("Total time: ", time.time() - start, "seconds")

```

```

Epoch 1/5
120/120 [=====] - 290s 2s/step - loss: 1.1276 - accuracy: 0.4862
- val_loss: 0.5640 - val_accuracy: 0.7650
Epoch 2/5
120/120 [=====] - 283s 2s/step - loss: 0.5131 - accuracy: 0.7983
- val_loss: 0.3830 - val_accuracy: 0.8514
Epoch 3/5
120/120 [=====] - 281s 2s/step - loss: 0.3902 - accuracy: 0.8560
- val_loss: 0.3422 - val_accuracy: 0.8744
Epoch 4/5
120/120 [=====] - 288s 2s/step - loss: 0.3277 - accuracy: 0.8822
- val_loss: 0.2951 - val_accuracy: 0.8902
Epoch 5/5
120/120 [=====] - 303s 3s/step - loss: 0.2972 - accuracy: 0.8944
- val_loss: 0.3023 - val_accuracy: 0.8842
Total time: 1462.7450199127197 seconds

```

```

In [ ]: #evaluation
_, train_acc = model.evaluate(x_train, y_train, verbose=0)
_, val_acc = model.evaluate(x_val, y_val, verbose=0)
_, test_acc = model.evaluate(x_test, y_test, verbose=0)
print('Train: %.3f, Val: %.3f,Test: %.3f' % (train_acc*100, val_acc*100,
test_acc*100))

#Time performance for test dataset
start_time = time.time()
y_pred = model.predict(x_test)
end_time = time.time()
time_taken = end_time - start_time
print("Testing time : ", time_taken)

```

```

Train: 89.415, Val: 88.417,Test: 88.040
Testing time : 13.466981887817383


```




```
In [ ]: plt.plot(CNN_VGG_history.history['accuracy'])
plt.plot(CNN_VGG_history.history['val_accuracy'])
plt.title('Model accuracy vs Epoch')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='lower right')
plt.show()

# summarize history for loss
plt.plot(CNN_VGG_history.history['loss'])
plt.plot(CNN_VGG_history.history['val_loss'])
plt.title('Model loss vs Epoch')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper right')
plt.show()

plt.plot(CNN_VGG_history.history['val_accuracy'])
plt.plot(CNN_VGG_history.history['val_loss'])
plt.title('Validation Accuracy vs Validation Loss based on Epochs')
plt.ylabel('Validation Accuracy, Test Loss')
plt.xlabel('Epochs')
plt.legend(['accuracy', 'loss'], loc='upper right')
plt.show()

plt.plot(CNN_VGG_history.history['accuracy'])
plt.plot(CNN_VGG_history.history['loss'])
plt.title('Training Accuracy vs Training Loss based on Epochs')
plt.ylabel('Train Accuracy, Train Loss')
plt.xlabel('Epochs')
plt.legend(['accuracy', 'loss'], loc='upper right')
plt.show()
```





## LSTM

Long Short Term Memory are special type of Recurrent Neural network (RNN) architecture i sable to learn long range dependencies. It has input and forget gate whose values are calculated based on previous cell state and previous hidden state.

```
In [ ]: row_hidden = 128  
col_hidden = 128
```

```

num_classes = 5
row, col, pixel = x_train.shape[1:]
x = Input(shape=(row, col, pixel))
encoded_rows = TimeDistributed(LSTM(row_hidden))(x)
encoded_columns = LSTM(col_hidden)(encoded_rows)
prediction = Dense(num_classes, activation='softmax')(encoded_columns)
model = Model(x, prediction)
model.compile(loss='categorical_crossentropy',
              optimizer='rmsprop',
              metrics=['accuracy'])

```

In [ ]: `model.summary()`

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 28, 28, 1)]	0
time_distributed (TimeDistributed)	(None, 28, 128)	66560
lstm_1 (LSTM)	(None, 128)	131584
dense (Dense)	(None, 5)	645
=====		

Total params: 198,789  
Trainable params: 198,789  
Non-trainable params: 0

In [ ]:

```

start = time.time()
model.fit(x_train, y_train,
          batch_size=BATCH_SIZE,
          epochs=EPOCH,
          verbose=1,
          validation_data=(x_val, y_val))
print("Total time: ", time.time() - start, "seconds")

```

```

Epoch 1/10
96/96 [=====] - 457s 5s/step - loss: 1.0675 - accuracy: 0.5205 -
val_loss: 0.7682 - val_accuracy: 0.6822
Epoch 2/10
96/96 [=====] - 447s 5s/step - loss: 0.7299 - accuracy: 0.6923 -
val_loss: 0.6500 - val_accuracy: 0.7227
Epoch 3/10
96/96 [=====] - 481s 5s/step - loss: 0.6108 - accuracy: 0.7523 -
val_loss: 0.5363 - val_accuracy: 0.7925
Epoch 4/10
96/96 [=====] - 489s 5s/step - loss: 0.5445 - accuracy: 0.7859 -
val_loss: 0.5266 - val_accuracy: 0.7949

```

```

Epoch 5/10
96/96 [=====] - 515s 5s/step - loss: 0.4983 - accuracy: 0.8077 -
val_loss: 0.4239 - val_accuracy: 0.8419
Epoch 6/10
96/96 [=====] - 512s 5s/step - loss: 0.4539 - accuracy: 0.8259 -
val_loss: 0.4301 - val_accuracy: 0.8392
Epoch 7/10
96/96 [=====] - 514s 5s/step - loss: 0.4164 - accuracy: 0.8430 -
val_loss: 0.3707 - val_accuracy: 0.8652
Epoch 8/10
96/96 [=====] - 539s 6s/step - loss: 0.3834 - accuracy: 0.8572 -
val_loss: 0.3659 - val_accuracy: 0.8696
Epoch 9/10
96/96 [=====] - 535s 6s/step - loss: 0.3596 - accuracy: 0.8654 -
val_loss: 0.3619 - val_accuracy: 0.8675
Epoch 10/10
96/96 [=====] - 535s 6s/step - loss: 0.3394 - accuracy: 0.8729 -
val_loss: 0.3818 - val_accuracy: 0.8575
Total time: 5064.516736030579 seconds

```

```

In [ ]: #evaluation
_, train_acc = model.evaluate(x_train, y_train, verbose=0)
_, val_acc = model.evaluate(x_val, y_val, verbose=0)
_, test_acc = model.evaluate(x_test, y_test, verbose=0)
print('Train: %.3f, Val: %.3f,Test: %.3f' % (train_acc*100, val_acc*100,
test_acc*100))

```

Train: 86.106, Val: 85.750,Test: 85.130

```

In [ ]: #Time performance for test dataset
start_time = time.time()
y_pred = model.predict(x_test)
end_time = time.time()
time_taken = end_time - start_time
print("Testing time : ", time_taken)

```

Testing time : 42.44548010826111

## Comparisons

Model	Training Accuracy	Validation Accuracy	Testing Accuracy	Model Training time	Model Testing time	Epochs trained	Optimizer	Loss function	Activation Function
CM1	100.00	94.492	93.960	2063.238160133362 s	3.376964807510376 s	50	Adam	Categorical Cross entropy	ReLU
CM2 (CNN1)	97.206	93.967	93.820	1222.489366531372 s	5.968332529067993 s	10	Adam	Categorical Cross entropy	ReLU
CNN2	93.448	92.683	92.420	322.5631422996521 s	1.8049075603485107 s	10	Adam	Categorical Cross entropy	ReLU
CNN3	75.352	75.400	75.210	322.4274568557739 s	1.7907054424285889 s	10	SGD	Categorical Cross entropy	ReLU
CNN4	94.356	93.258	92.960	322.5065915584564 s	2.2175590991973877 s	10	Adam	Categorical Cross entropy	Leaky ReLU
CNN_VGG	89.415	88.417	88.040	1462.745019912719 s	13.466981887817383 s	5	Adam	Categorical Cross entropy	ReLU
CNN_LSTM	86.106	85.750	85.130	5064.516736030579 s	42.44548010826111 s	10	Adam	Categorical Cross entropy	ReLU

## Observations:

When trained with same Epoch =10 and batch size of 500 images , the models performed similar except CNN3(SGD optimizer) and CNN\_VGG and LSTM architectures. This is possibly as SGD takes longer to converge but generalizes better(argued in various papers). LSTM and CNN\_VGG are complex models with greater number of parameters to learn. This is also apparent through the runtime performance for these models. The plots for model accuracy vs epochs and model loss vs

epoch are plotted above and it is observed that in complex models it does not reach its optimal training accuracy and loss did not start increasing whereas in simpler models, models trained much faster and tend to overfit (as apparent by increase in loss function).

## CM4 : Using your own encoding

```
In [ ]: model = Sequential(name="CNN_2")

model.add(Conv2D(32, kernel_size=(3, 3),
                strides=(1, 1),
                activation='relu',
                input_shape=(28,28,1), padding = 'same', name = 'Conv1'))
model.add(MaxPooling2D(pool_size=(2, 2), name = 'MaxPool1'))
#model.add(Dropout(0.3))

model.add(Conv2D(32, (3, 3), activation='relu', name = 'Conv2'))
model.add(MaxPooling2D(pool_size=(2, 2), name = 'MaxPool2'))
model.add(Dropout(0.3))

model.add(Flatten(name = 'Flatten'))

model.add(Dense(256, activation='relu', name = 'Dense1'))
model.add(Dropout(0.5))

model.add(Dense(64, activation='relu', name = 'Dense2'))

model.add(Dense(5, activation='softmax', name = 'Softmax'))

model.compile(loss='categorical_crossentropy', optimizer='adam',
              metrics=['accuracy'])
```

```
In [ ]: CNN_2_history = model.fit(x_train, y_train,
                                batch_size=BATCH_SIZE,
                                epochs=3,
                                shuffle = True,
                                validation_data=(x_val, y_val))
```

Epoch 1/3

96/96 [=====] - 39s 388ms/step - loss: 0.6579 - accuracy: 0.7325  
- val\_loss: 0.3782 - val\_accuracy: 0.8569

Epoch 2/3

96/96 [=====] - 32s 338ms/step - loss: 0.3884 - accuracy: 0.8553  
- val\_loss: 0.3047 - val\_accuracy: 0.8897

Epoch 3/3

```
96/96 [=====] - 32s 337ms/step - loss: 0.3304 - accuracy: 0.8783  
- val_loss: 0.2699 - val_accuracy: 0.8997
```

## Intermediate Layer Model

```
In [ ]: dense_layer_output = [layer.output for layer in model.layers if layer.name ==  
    'Dense2']  
  
extractor = Model(inputs=model.inputs, outputs=dense_layer_output)  
  
features = extractor(x_test)
```

```
In [ ]: print(type(features))  
print(len(features))  
print(len(x_test))
```

```
<class 'tensorflow.python.framework.ops.EagerTensor'>  
10000  
10000
```

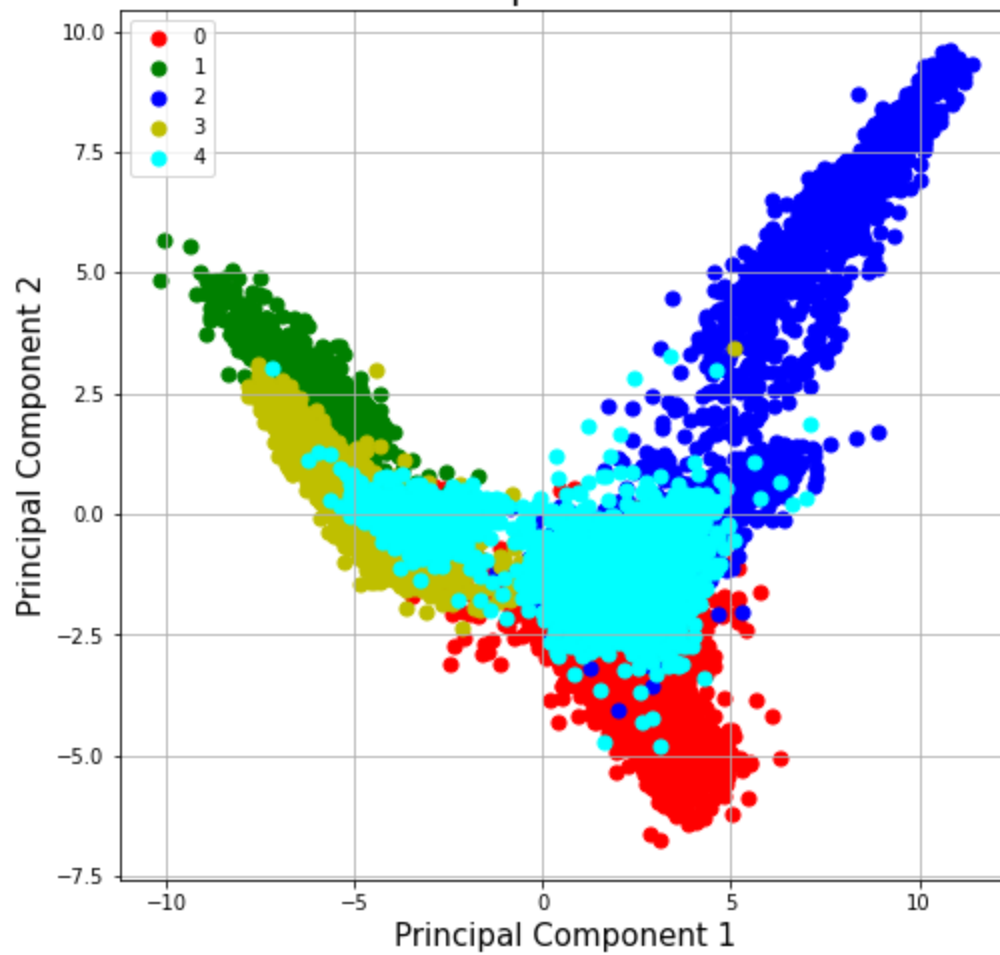
## PCA on extracted features

```
In [ ]: pca = PCA(n_components=n_components, random_state=random_state)  
pca.fit(features)  
x = pca.transform(features)
```

```
In [ ]: principalDf = pd.DataFrame(data = x  
    , columns = ['principal component 1', 'principal component 2'])  
finalDf = pd.concat([principalDf, df_y_test[['0']], axis = 1)
```


```
In [ ]: fig = plt.figure(figsize = (8,8))  
ax = fig.add_subplot(1,1,1)  
ax.set_xlabel('Principal Component 1', fontsize = 15)  
ax.set_ylabel('Principal Component 2', fontsize = 15)  
ax.set_title('2 component PCA', fontsize = 20)  
targets = classes  
colors = ['r', 'g', 'b', 'y', 'cyan']  
for target, color in zip(targets, colors):  
    indicesToKeep = finalDf['0'] == target  
    ax.scatter(finalDf.loc[indicesToKeep, 'principal component 1']  
        , finalDf.loc[indicesToKeep, 'principal component 2']  
        , c = color  
        , s = 50)  
  
ax.legend(targets)  
ax.grid()
```

# 2 component PCA



```
In [ ]: dimReducedDataFrame = pd.DataFrame(x)
dimReducedDataFrame = dimReducedDataFrame.rename(columns = { 0: 'V1', 1 :
'V2' })
dimReducedDataFrame['Category'] = df_y_test['0']
plt.figure(figsize = (10, 5))
sb.scatterplot(data = dimReducedDataFrame, x = 'V1', y = 'V2', hue =
'Category')
plt.grid(True)
plt.show()
```





## KMeans


```
In [ ]: kmeans = KMeans(n_clusters=5, random_state=22)
kmeans.fit(features)
```

```
Out[ ]: KMeans(n_clusters=5, random_state=22)
```

```
In [ ]: #kmeans = KMeans(n_clusters=5, random_state=random_state)
#z=kmeans.fit_transform(features)
```

```
In [ ]: label = kmeans.fit_predict(features)
u_labels = np.unique(label)
fig = plt.figure(figsize = (8,8))
ax = fig.add_subplot(1,1,1)
ax.set_xlabel('Component 1', fontsize = 15)
ax.set_ylabel('Component 2', fontsize = 15)
ax.set_title('K-MEANS Clustering', fontsize = 20)
fig = plt.figure(figsize = (8,8))
for i in u_labels:
    ax.scatter(x[label == i , 0] , x[label == i , 1] , label = i)
ax.legend()
ax.grid()
```

## K-MEANS Clustering



<Figure size 576x576 with 0 Axes>

```
In [ ]: #unique, counts = np.unique(label, return_counts=True)
        #print(unique, counts)
```

```
[0 1 2 3 4] [4951 2113  848 1022 1066]
```

```
In [ ]: #df_y_test['0'].value_counts()
```

```
Out[ ]: 4      3000
        2      2000
        0      2000
        3      2000
        1      1000
        Name: 0, dtype: int64
```


## DBSCAN

### Finding optimal value of epsilon

```
In [ ]: neigh = NearestNeighbors(n_neighbors=2)
        nbrs = neigh.fit(features)
        distances, indices = nbrs.kneighbors(features)
```

```
In [ ]: distances = np.sort(distances, axis=0)
        distances = distances[:,1]
```

```
plt.figure(figsize=(8,6))
plt.plot(distances)
plt.title('K-distance Graph',fontsize=20)
plt.xlabel('Data Points sorted by distance',fontsize=14)
plt.ylabel('Epsilon',fontsize=14)
plt.show()
```



## DBSCAN using extracted features

```
In [ ]: dbscan_opt=DBSCAN(eps=1.6,min_samples=60)
#d = dbscan_opt.fit_predict(features)
d = dbscan_opt.fit(features)
```

```
In [ ]: principalDf['DBSCAN_opt_labels'] = dbscan_opt.labels_
print(dbscan_opt.labels_)
print(np.unique(dbscan_opt.labels_))
print(principalDf['DBSCAN_opt_labels'].value_counts())
```


```
[ 0 -1  1 ... -1  1  0]
[-1  0  1  2  3  4]
2    4626
0    1975
-1   1281
1     839
3     699
4     580
```

Name: DBSCAN\_opt\_labels, dtype: int64

```
In [ ]: import matplotlib
plt.figure(figsize=(10,10))
```

```
plt.scatter(principalDf['principal component 1'],principalDf['principal
component 2'],c=principalDf['DBSCAN_opt_labels'],
           cmap=matplotlib.colors.ListedColormap(colors),s=15)
plt.title('DBSCAN Clustering',fontsize=20)
plt.xlabel('Feature 1',fontsize=14)
plt.ylabel('Feature 2',fontsize=14)
plt.legend()
plt.show()
```

No handles with labels found to put in legend.



## T-SNE

```
In [ ]: tsne = TSNE(n_components=n_components, random_state=random_state)#verbose=1
z = tsne.fit_transform(features)
```

/usr/local/lib/python3.7/dist-packages/sklearn/manifold/\_t\_sne.py:783: FutureWarning: The default initialization in TSNE will change from 'random' to 'pca' in 1.2.

FutureWarning,


/usr/local/lib/python3.7/dist-packages/sklearn/manifold/\_t\_sne.py:793: FutureWarning: The

default learning rate in TSNE will change from 200.0 to 'auto' in 1.2.  
FutureWarning,

```
In [ ]: tsneDf = pd.DataFrame(data = z
                             , columns = ['component 1', 'component 2'])
finalDf = pd.concat([tsneDf, df_y_test[['0']], axis = 1)
```

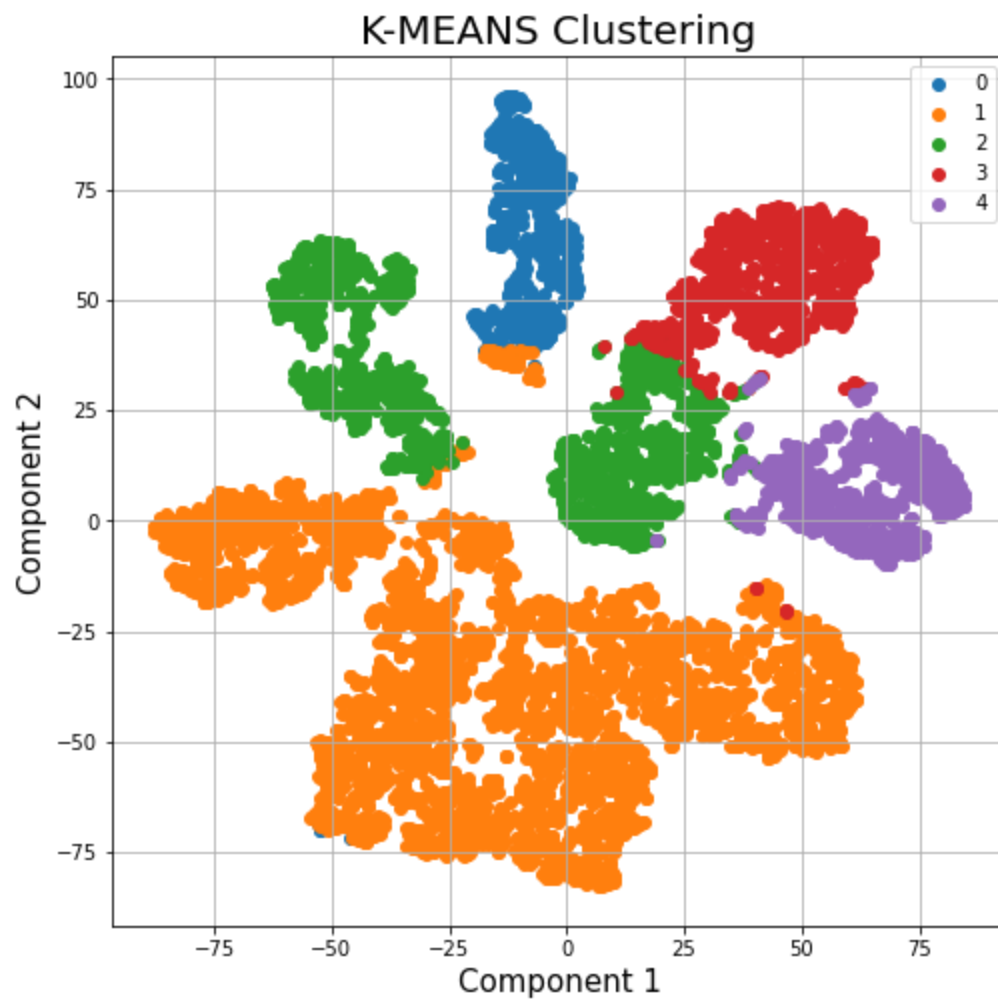
```
In [ ]: sb.scatterplot(x="component 1", y="component 2", hue=finalDf['0'],
                      palette=sb.color_palette("hls", 5),
                      data=finalDf).set(title="T-SNE projection")
```

Out[ ]: [Text(0.5, 1.0, 'T-SNE projection')]



## K-Means Clustering

```
In [ ]: kmeans = KMeans(n_clusters=5, random_state=22)
kmeans.fit(features)
label = kmeans.fit_predict(features)
u_labels = np.unique(label)
fig = plt.figure(figsize = (8,8))
ax = fig.add_subplot(1,1,1)
ax.set_xlabel('Component 1', fontsize = 15)
ax.set_ylabel('Component 2', fontsize = 15)
ax.set_title('K-MEANS Clustering', fontsize = 20)
fig = plt.figure(figsize = (8,8))
for i in u_labels:
    ax.scatter(z[label == i , 0] , z[label == i , 1] , label = i)
ax.legend()
ax.grid()
```




<Figure size 576x576 with 0 Axes>

## DBSCAN Clustering

```
In [ ]: neigh = NearestNeighbors(n_neighbors=2)
nbrs = neigh.fit(features)
distances, indices = nbrs.kneighbors(features)
distances = np.sort(distances, axis=0)
distances = distances[:,1]
plt.figure(figsize=(8,6))
plt.plot(distances)
plt.title('K-distance Graph',fontsize=20)
plt.xlabel('Data Points sorted by distance',fontsize=14)
plt.ylabel('Epsilon',fontsize=14)
plt.show()
```

## K-distance Graph

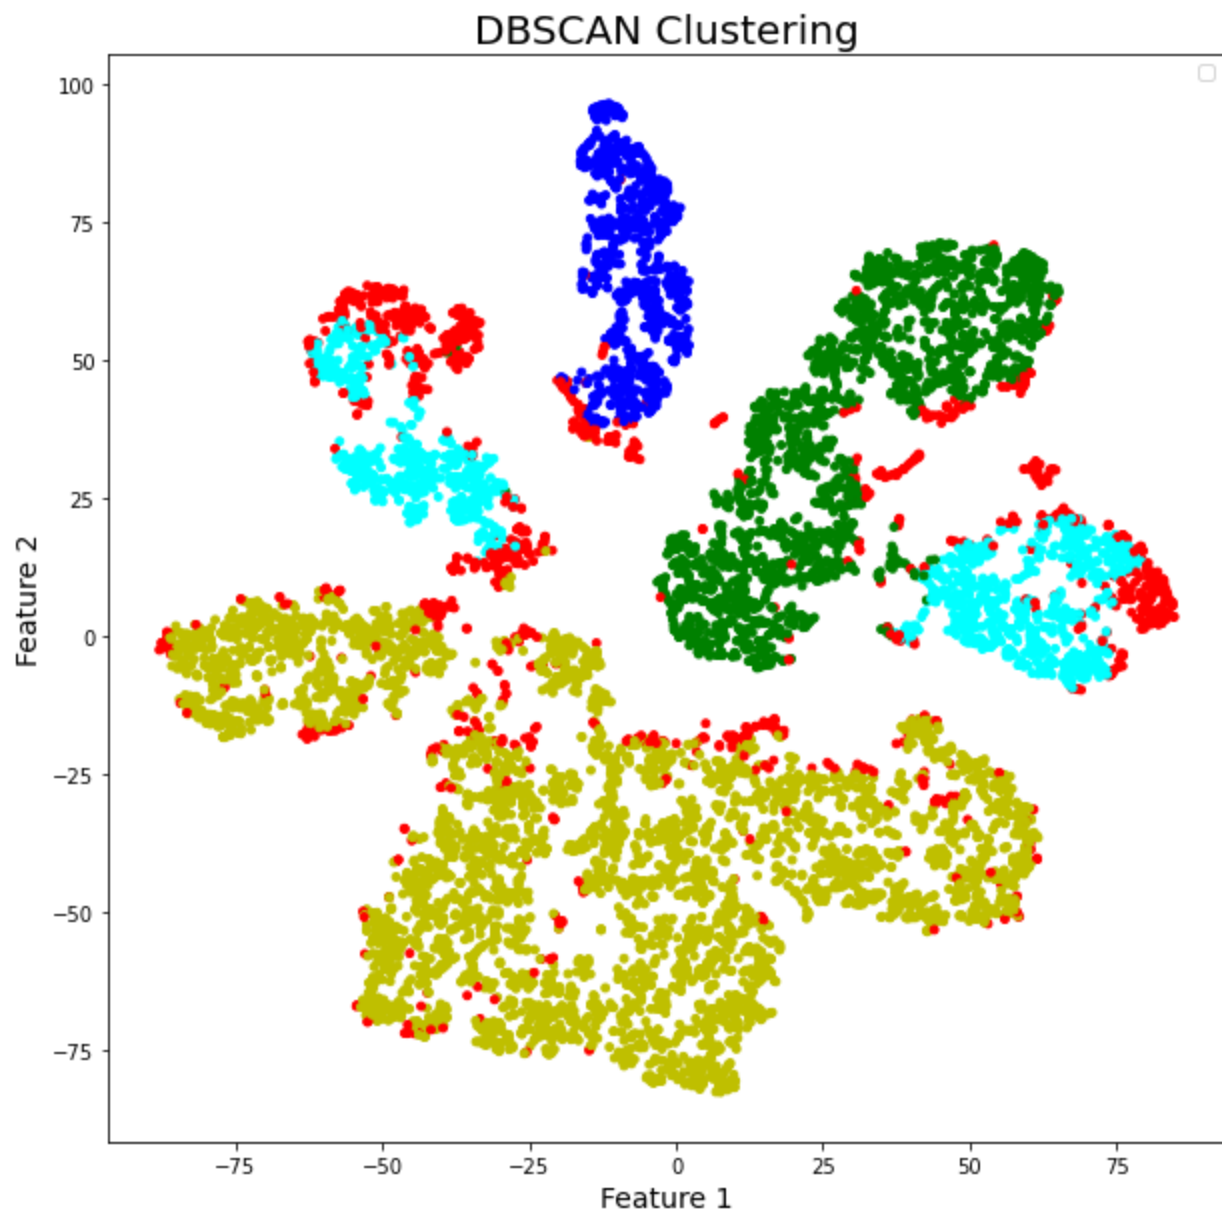


```
In [ ]: dbscan_opt=DBSCAN(eps=1.6,min_samples=60)
#d = dbscan_opt.fit_predict(features)
d = dbscan_opt.fit(features)
finalDf['DBSCAN_opt_labels'] = dbscan_opt.labels_
print(dbscan_opt.labels_)
print(np.unique(dbscan_opt.labels_))
print(finalDf['DBSCAN_opt_labels'].value_counts())
```

```
[ 0 -1  1 ... -1  1  0]
[-1  0  1  2  3  4]
 2    4626
 0    1975
-1    1281
 1     839
 3     699
 4     580
Name: DBSCAN_opt_labels, dtype: int64
```

```
In [ ]: plt.figure(figsize=(10,10))
plt.scatter(finalDf['component 1'],finalDf['component
2'],c=finalDf['DBSCAN_opt_labels'],
           cmap=matplotlib.colors.ListedColormap(colors),s=15)
plt.title('DBSCAN Clustering',fontsize=20)
plt.xlabel('Feature 1',fontsize=14)
plt.ylabel('Feature 2',fontsize=14)
plt.legend()
plt.show()
```

No handles with labels found to put in legend.



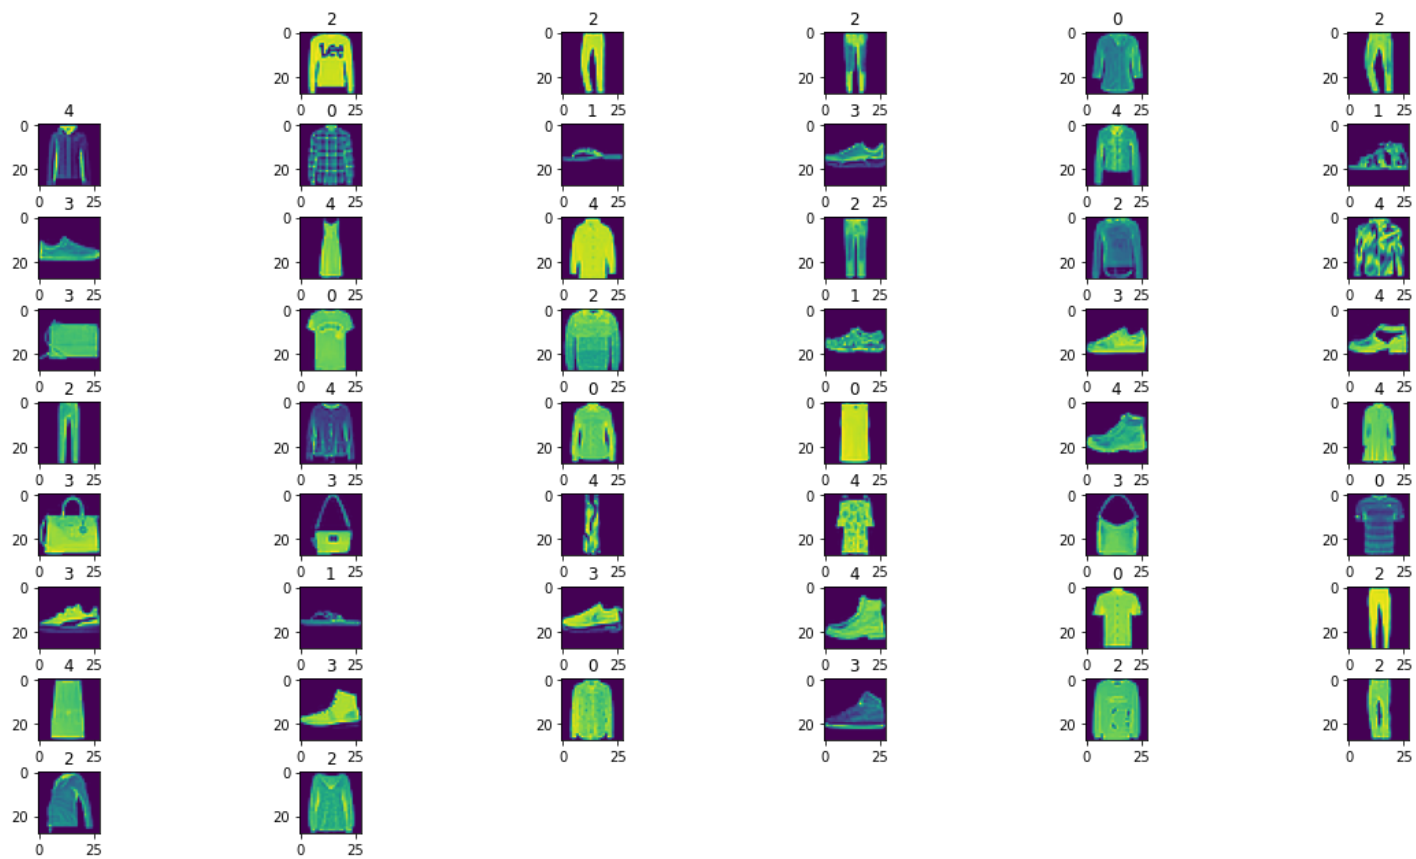
## Decoding mystery labels

In [24]:

```
from PIL import Image
rows=10
plt.figure(figsize=(20, 12))
plt.subplots_adjust(hspace=0.5)
plt.suptitle("images", fontsize=18, y=0.95)

for i in range(1,50):
    #print(i)
    #print(df_x_test[:i].shape)
    img=df_x_test.iloc[i].values.reshape((28,28,1))
    pil_img = tf.keras.preprocessing.image.array_to_img(img)
    plt.subplot(rows,6,i+1)
    plt.title(df_y_test['0'].iloc[i])
    plt.imshow(pil_img)
    #plt.show()
```





## Mystery Labels Decoded

We plotted images with their true labels from test dataset and with predicted labels using KMeans on features encoding applied on test dataset. These are the major classes/ labels that are associated with category, although broader category (Summer wear, Winterwear or clothing material ) is hard to define because of certain labels having mixed classes, we define grouped labels on the granular/ initial fashion mnist dataset labels that could have been represented by our mystery labels.

- Label 0 : Tshirt/Top and Shirt
- Label 1 : Sandals
- Label 2 : Trouser , Coat and Pullover
- Label 3 : Sneakers and Handbags
- Label 4 : Dress and Ankle Boots

## References

[1] Gulli, A., & Pal, S. (2017). Deep learning with Keras. Packt Publishing Ltd. [2] Retrieved from [https://colab.research.google.com/github/markwest1972/CNN-Example-Google-Colaboratory/blob/master/CNN\\_Fashion\\_MNIST\\_Example.ipynb](https://colab.research.google.com/github/markwest1972/CNN-Example-Google-Colaboratory/blob/master/CNN_Fashion_MNIST_Example.ipynb) [3] Compare the effect of different scalers on data with outliers — scikit-learn 0.22.1 documentation. (n.d.). Retrieved January, 2022, from [https://scikitlearn.org/stable/auto\\_examples/preprocessing/plot\\_all\\_scaling.html#sphx-glr-auto-examples-preprocessing-plot-all-scaling-py](https://scikitlearn.org/stable/auto_examples/preprocessing/plot_all_scaling.html#sphx-glr-auto-examples-preprocessing-plot-all-scaling-py) contributed by Scikit-learn: Machine Learning in Python, Pedregosa et al., JMLR 12, pp. 2825-2830, 2011. [4] Discussed KNeighbor Classifiers — scikit-learn 0.22.1 documentation. (n.d.). Retrieved January 27, 2022, from <https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html> contributed by Scikit-learn: Machine Learning in Python, Pedregosa et al., JMLR 12, pp. 2825-2830, 2011. [5] Discussed Normalization techniques and parameters — scikit-learn 0.22.1 documentation. (n.d.). Retrieved January 27, 2022, from

<https://scikit-learn.org/stable/modules/preprocessing.html#normalization> contributed by Scikit-learn: Machine Learning in Python, Pedregosa et al., JMLR 12, pp. 2825-2830, 2011. [6] Discussed different distance metrics — scikit-learn 0.22.1 documentation. (n.d.). Retrieved January 27, 2022, from <https://scikit-learn.org/stable/modules/generated/sklearn.metrics.DistanceMetric.html#sklearn.metrics.DistanceMetric>

[7] Discussed axes, plotting and figures in python — matplotlib.org documentation. (n.d.). Retrieved January 27, 2022, from [https://matplotlib.org/stable/api/legend\\_api.html?highlight=legends](https://matplotlib.org/stable/api/legend_api.html?highlight=legends) contributed by J. D. Hunter, "Matplotlib: A 2D Graphics Environment", Computing in Science & Engineering, vol. 9, no. 3, pp. 90-95, 2007. [8] Discussed plotting of dataframe columns in python — pandas.org documentation. (n.d.). Retrieved January 27, 2022, from <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.plot.scatter.html> contributed by McKinney, W., & others. (2010). Data structures for statistical computing in python. In Proceedings of the 9th Python in Science Conference (Vol. 445, pp. 51–56). [9] Retrieved from <https://github.com/pytorch/vision/blob/6db1569c89094cf23f3bc41f79275c45e9fcb3f3/torchvision/models/vgg.py#L24>

[10] Brownlee, J. (2019). Ordinal and One-Hot Encodings for Categorical Data. Retrieved from <https://machinelearningmastery.com/> [11] Katara, V. (2020). Dimensionality Reduction — PCA vs LDA vs t-SNE. Retrieved from <https://medium.com/analytics-vidhya/dimensionality-reduction-pca-vs-lda-vs-t-sne-681636bc686>

[12] shakhadri313, (2021). Build VGG Net from Scratch with Python!. Retrieved from <https://www.analyticsvidhya.com/blog/2021/06/build-vgg-net-from-scratch-with-python/> [13] Brownlee, J. (2019). how-to-visualize-filters-and-feature-maps-in-convolutional-neural-networks from <https://machinelearningmastery.com/how-to-visualize-filters-and-feature-maps-in-convolutional-neural-networks/>

[14] Harsh (2020) Classifying Fashion\_Mnist dataset with Convolutional Neural Nets Retrieved from <https://medium.com/swlh/classifying-fashion-mnist-dataset-with-convolutional-neural-nets-dd092d755164>